







RAAZ DWIVEDI

 Harvard/MIT  rzrsk.github.io  raaz@seas.harvard.edu  raaz@mit.edu    

ACADEMIC APPOINTMENTS

FODSI Postdoctoral Fellow 2021—
Harvard University, School of Engineering, and Applied Sciences, Boston, USA
Massachusetts Institute of Technology, Department of EECS, Cambridge, USA
• Advisors: *Prof. Susan Murphy & Prof. Devavrat Shah*

EDUCATION

Ph.D., Electrical Engineering, and Computer Sciences 2015—2021
University of California, Berkeley, USA
• Advisors: *Prof. Martin Wainwright & Prof. Bin Yu*
• Thesis title: *Principled statistical approaches for sampling, and inference in high dimensions*
• Thesis committee members: *Prof. David Aldous & Prof. Peter Bartlett*
B. Tech., Honors, Electrical Engineering 2010—2014
Indian Institute of Technology, Bombay, India
• Advisor: *Prof. Vivek Borkar*
• Graduated with Institute Rank 1, and Minors in Mathematics

RESEARCH INTERESTS

Theoretical, and applied aspects of statistical machine learning, and data science with a focus on causal inference, reinforcement learning, and theory of MCMC methods

ACHIEVEMENTS & AWARDS

Institute of Mathematical Statistics (IMS) New Researcher Travel Award, London 2022
Best Presentation Award, Machine Learning, and Statistics Session, Laboratory of Information, and Decision Systems (LIDS) Student Conference, MIT 2022
Best Student Paper Award, Sections on Statistical Computing, and Statistical Graphics, American Statistical Association (ASA) 2022
Postdoctoral Fellowship, Foundations of Data Science Institute (FODSI) 2021
Outstanding Graduate Student Instructor Award, UC Berkeley 2020
Berkeley Student Travel Grant, Jerusalem Joint Statistical Event, Israel 2018
Student Travel Award, NeurIPS 2018, Montreal Canada 2018
Leibniz Graduate Students Travel Award, Oberwolfach, Germany 2017
Student Travel Award, SAMSI QMC Workshop 2017, Raleigh-Durham 2017
Berkeley Fellowship, the most prestigious fellowship for incoming students 2015
President of India Gold Medal, IIT Bombay, for the highest GPA in the institute 2014
Institute Silver Medal, IIT Bombay, for the highest Honors GPA (EE department) 2014
Best B. Tech. Project (undergraduate thesis) Award, IIT Bombay 2014
Manjula Bagmal Parikh Foundation Trust Prize, IIT Bombay 2014
Dilip R. Limaye Academic Excellence Award, IIT Bombay 2014
Prof. K. C. Mukherji Award, IIT Bombay 2014
Urvish Medh Memorial Prize for excellent academic performance in the EE department in first through senior year, IIT Bombay 2011—2014

Institute Academic Excellence Prize for excellent institute wide academic performance in first through senior year, IIT Bombay 2011—2014

Aditya Choubey Prize for the best academic performance in the institute across the first year students, IIT Bombay 2011

All India Rank 10 (amongst half a million), IIT Joint Entrance Exam (IIT-JEE) 2010

All India Rank 46 (amongst a million), All India Engineering Entrance Exam 2010

WORK EXPERIENCE **Microsoft Research**, Research Intern (with Lester Mackey), New England, USA Summer 2019
Mist Systems (Juniper Networks), Data Science Intern, Cupertino, USA Summer 2017
WorldQuant Research, Senior Quantitative Researcher, Mumbai, India Jul 2014—Jul 2015
Stanford University, Research intern (with Prof. Balaji Prabhakar), USA Summer 2013
Ivy Mobility, Data Science Intern, Chennai, India Winter 2012

JOURNAL PUBLICATIONS

(★ denotes equal contribution, and † denotes alphabetical ordering; title is hyperlinked to the online pdf of the paper)

- J1. **Raaz Dwivedi**★, Yan Shuo Tan★, Briton Park, Mian Wei, Kevin Horgan, David Madigan, and Bin Yu, “Stable discovery of interpretable subgroups via calibration in causal studies”, *International Statistical Review (ISR)*, 2020.
- J2. Nick Altieri†, Rebecca L. Barter, James Duncan, **Raaz Dwivedi**, Karl Kumbier, Xiao Li, Robert Netzer, Briton Park, Chandan Singh, Yan Shuo Tan, Tiffany Tang, Yu Wang, Chao Zhang, and Bin Yu, “Curating a COVID-19 data repository and forecasting county-level death counts in the United States”, *Harvard Data Science Review (HDSR)*, 2020.
- J3. **Raaz Dwivedi**★, Nhat Ho★, Koulik Khamaru★, Martin J. Wainwright, Michael I. Jordan, and Bin Yu, “Singularity, misspecification, and the convergence rate of EM”, *Annals of Statistics (AoS)*, 2020.
- J4. Yuansi Chen, **Raaz Dwivedi**, Martin J. Wainwright, and Bin Yu, “Fast mixing of Metropolized Hamiltonian Monte Carlo: Benefits of multi-step gradients”, *Journal of Machine Learning Research (JMLR)*, 2020.
- J5. **Raaz Dwivedi**★, Yuansi Chen★, Martin J. Wainwright, and Bin Yu, “Log-concave sampling: Metropolis-Hastings algorithms are fast”, *Journal of Machine Learning Research (JMLR)*, 2019.
- J6. **Raaz Dwivedi**†, Ohad N. Feldheim, Ori Gurel-Gurevich, and Aaditya Ramdas, “The power of online thinning in reducing discrepancy”, *Probability Theory, and Related Fields (PTRF)*, 2019.
- J7. Yuansi Chen★, **Raaz Dwivedi**★, Martin J. Wainwright, and Bin Yu, “Fast MCMC sampling algorithms on polytopes”, *Journal of Machine Learning Research (JMLR)*, 2018.
- J8. Vivek Borkar†, **Raaz Dwivedi**, and Neeraja Sahasrabudhe, “Gaussian approximations in high dimensional estimation”, *Systems & Control Letters*, 2016.

CONFERENCE PUBLICATIONS

- C1. **Raaz Dwivedi**, and Lester Mackey, “Generalized kernel thinning”, To appear in *International Conference on Learning Representations (ICLR)*, 2022.
- C2. Abhishek Shetty, **Raaz Dwivedi**, and Lester Mackey, “Distribution compression in near-linear time”, To appear in *International Conference on Learning Representations (ICLR)*, 2022.
- C3. **Raaz Dwivedi**, and Lester Mackey, “Kernel thinning”, Extended abstract in *Conference on Learning Theory (COLT)*, 2021.
- C4. **Raaz Dwivedi**★, Nhat Ho★, Koulik Khamaru★, Martin J. Wainwright, Michael I. Jordan, and Bin Yu, “Sharp analysis of Expectation-Maximization for weakly identifiable models”, *The 23rd International Conference on Artificial Intelligence, and Statistics (AISTATS)*, 2020.

- C5. **Raaz Dwivedi***, Nhat Ho*, Koulik Khamaru*, Martin J. Wainwright, and Michael I. Jordan, “Theoretical guarantees for EM under misspecified Gaussian mixture models”, *Advances in Neural Information Processing Systems (NeurIPS)*, Montréal, 2018.
- C6. **Raaz Dwivedi***, Yuansi Chen*, Martin J. Wainwright, and Bin Yu, “Log-concave sampling: Metropolis-Hastings algorithms are fast”, Extended abstract in *Conference on Learning Theory (COLT)*, Stockholm, 2018.
- C7. Yuansi Chen*, **Raaz Dwivedi***, Martin J. Wainwright, and Bin Yu, “Vaidya walk: A sampling algorithm based on the volumetric barrier”, *Communication, Control, and Computing (Allerton)*, 55th Annual Allerton Conference, 2017.
- C8. **Raaz Dwivedi**, and Vivek Borkar, “Removing sampling bias in networked stochastic approximation”, *International Conference on Signal Processing, and Communications (SPCOM)*, Bangalore, 2014.

PRE-PRINTS

- P1. **Raaz Dwivedi**, Susan Murphy, and Devavrat Shah, “Counterfactual inference in sequential experimental design”, *arXiv preprint 2022*.
- P2. **Raaz Dwivedi***, Chandan Singh*, Bin Yu, and Martin J. Wainwright, “Revisiting minimum description length complexity in overparameterized models”, *arXiv preprint 2021 (in journal submission)*.
- P3. Nhat Ho*, Koulik Khamaru*, **Raaz Dwivedi***, Martin J. Wainwright, Michael I. Jordan, and Bin Yu, “Instability, computational efficiency, and statistical accuracy”, *arXiv preprint 2021 (in journal submission)*.

SOFTWARES & METHODOLOGIES

- P1. **Raaz Dwivedi**, Lester Mackey, Python package “Kernel Thinning”, available at <https://github.com/microsoft/goodpoints>.
- P2. Abhishek Shetty*, **Raaz Dwivedi***, and Lester Mackey, Python package “Compress++”, available at <https://github.com/microsoft/goodpoints>.
- P3. **Raaz Dwivedi***, Yan Shuo Tan*, Briton Park, Mian Wei, Kevin Horgan, David Madigan, and Bin Yu, GitHub Repo “StaDISC” available at <https://github.com/Yu-Group/stadisc>.
- P4. Yuansi Chen*, **Raaz Dwivedi***, Martin Wainwright, and Bin Yu, Python package (C++ implementation) “Vaidya and John walks”, available at <https://github.com/yuachen/polytopewalk>.

INVITED RESEARCH TALKS

- IT1. Revisiting minimum description length complexity in overparameterized models. *Algorithmic Information Theory & Machine Learning Symp.*, Alan Turing Institute, London. July 2022
- IT2. Counterfactual inference in sequential experimental design. *Institute of Mathematical Statistics (IMS) Annual Meeting, Statistical Machine Learning Session*, London. June 2022
- IT3. Counterfactual inference in sequential experimental design. *Causal Inference Reading Group*, University of Cambridge. June 2022
- IT4. Near-optimal compression in near-linear time. *Sympo. on Kernel methods for numerical integration, SIAM Conference on Uncertainty Quantification*, Atlanta. Apr 2022
- IT5. Near-optimal compression in near-linear time. *Stable, Generalizable, & Transferable Statistical Learning Workshop*, Mathematical Sciences Research Institute, Berkeley. Mar 2022
- IT6. Counterfactual inference in sequential experimental design. *Learning from Interventions Workshop*, Simons Institute, Berkeley. Feb 2022
- IT7. Imputation using nearest neighbors for adaptively collected data. *Foundations of Data Science Institute (FODSI) Retreat*. Jan 2022
- IT8. Revisiting minimum description length complexity in overparameterized models. *Collaborations on the Theoretical Foundations of Deep Learning*. Nov 2021

- IT9. Non-asymptotic guarantees for MCMC, and Kernel thinning. *Finale Doshi-Velez Group Meeting, Harvard University.* Oct 2021
- IT10. Kernel thinning. *Data-Centric Eng. Group (DCE), Alan Turing Institute.* Sep 2021
- IT11. Subgroup discovery in randomized experiments & Markov chain Monte Carlo sampling. *Research Seminar, USC Marshall School of Business.* Feb 2021
- IT12. Subgroup discovery in randomized experiments & Markov chain Monte Carlo sampling. *Statistics Seminar, University of Toronto.* Feb 2021
- IT13. Subgroup discovery in randomized experiments & Markov chain Monte Carlo sampling. *MINDS Symposium on the Foundations of Data Science, JHU.* Feb 2021
- IT14. Subgroup discovery in randomized experiments & Markov chain Monte Carlo sampling. *Devavrat Shah's & Susan Murphy's Labs, MIT & Harvard University.* Feb 2021
- IT15. Subgroup discovery in randomized experiments & Markov chain Monte Carlo sampling. *Research Seminar, Microsoft Research New England.* Jan 2021
- IT16. Non-asymptotic guarantees for Markov chain Monte Carlo. *Flatiron Institute.* Jan 2021
- IT17. Subgroup discovery in randomized experiments & Markov chain Monte Carlo sampling. *Statistics Seminar, University of Washington.* Jan 2021
- IT18. Subgroup discovery in randomized experiments & Markov chain Monte Carlo sampling. *Operations Research, and Statistics Group Seminar, MIT Sloan.* Jan 2021
- IT19. New perspectives on old Problems in causal inference, and MCMC sampling. *Statistics Seminar, UC Irvine.* Jan 2021
- IT20. StaDISC: Stable discovery of interpretable subgroups via calibration. *Young Data Scientist Research Seminar, ETH Zurich.* Sep 2020
- IT21. Veridical data science, and the PCS framework. *ASA Annual Symposium on Data Science, and Statistics (SDSS).* Jun 2020
- IT22. Statistics meets optimization: Two vignettes on the Intersection. *Department of Mathematics, and Statistics, IIT Kanpur.* Jan 2020
- IT23. Singularity, misspecification, and the convergence rate of Expectation-Maximization. *Fall Western Sectional Meeting of the AMS, UC Riverside.* Nov 2019
- IT24. Power of gradients, and accept-reject step in MCMC algorithms. *BIDS Statistics, and Machine Learning Forum, UC Berkeley.* Mar 2019
- IT25. Theoretical guarantees for MCMC algorithms, *EE Dept, IIT Bombay.* Jan 2018
- IT26. Theoretical guarantees for MCMC algorithms, *STCS Seminar, TIFR Bombay.* Jan 2018

CONTRIBUTED
& OTHER
RESEARCH
TALKS

- CT1. Counterfactual inference in sequential experimental design. *MIT Statistics, and Data Science Conference (SDSCon) 2022, MIT.* April 2022
- CT2. Counterfactual inference in sequential experimental design. *Econometrics Lunch, MIT.* Mar 2022
- CT3. Near-optimal compression in near-linear time. *27th Annual Laboratory for Information & Decision Systems (LIDS) Student Conference, MIT.* Jan 2022
- CT4. Kernel thinning. *Stat 300, Harvard University.* Sep 2021
- CT5. Kernel thinning. *Monte Carlo Methods & Applications (MCM).* Sep 2021
- CT6. Kernel thinning. *Int. Soc. for Bayesian Analysis (ISBA) World Meeting 2021.* Aug 2021
- CT7. Kernel thinning. *The Bayesian Young Statisticians Meeting (BAYSM) 2021.* Aug 2021
- CT8. Kernel thinning. *Conference on Learning Theory (COLT).* Aug 2021
- CT9. Kernel thinning. *Subset Selection in Machine Learning Workshop, International Conference on Machine Learning (Subset ML, ICML). (Spotlight talk)* Jul 2021

CONTRIBUTED
POSTER
PRESENTATIONS

- CT10. Revisiting complexity, and the bias-variance tradeoff: Using minimum description length. *Workshop on the Theory of Overparameterized Machine Learning (TOPML)*. Apr 2021
- CT11. Log-concave sampling: Metropolis Hastings algorithms are fast. *Jerusalem Joint Statistical Event, Israel*. Dec 2018
- CP1. Counterfactual inference in sequential experimental design. *Synthetic Control Methods Workshop, Data X, 2022, Princeton University*. Jun 2022
- CP2. Counterfactual inference in sequential experimental design. *American Causal Inference Conference, 2022, UC Berkeley*. May 2022
- CP3. Counterfactual inference in sequential experimental design. *SMaSH: Symposium for Mathematical Sciences at Harvard, 2022, Harvard University*. May 2022
- CP4. Counterfactual inference in sequential experimental design. *MIT Statistics, and Data Science Conference (SDSCon) 2022, MIT*. April 2022
- CP5. Generalized kernel thinning. *Advances in Approx. Bayesian Inference (AABI)*. Feb 2022
- CP6. Revisiting complexity, and the bias-variance tradeoff: Using minimum description length. *North American School of Information Theory (NASIT)*. Jun 2021
- CP7. Log-concave sampling: Metropolis Hastings algorithms are fast. *Conference on Learning Theory (COLT) 2018, Stockholm, Sweden*. Dec 2018
- CP8. Theoretical guarantees for EM under misspecified Gaussian mixture models. *Neural Information Processing Systems (NeurIPS) 2018, Montréal, Canada*. Dec 2018
- CP9. The power to two choices in reducing discrepancy, *SAMSI QMC Opening Workshop, Raleigh-Durham, Duke University*. Aug 2017

TEACHING
EXPERIENCE

Teaching Fellow, Harvard University 2022–

- *Sequential Decision Making*: STAT 234, Spring 2022, Taught by Prof. Susan Murphy. Besides usual duties, gave guest lectures, and mentored several half semester-long class projects.

Graduate Student Instructor, UC Berkeley 2018–2019

- *Modern Statistical Prediction, and Machine Learning*: STAT 154, Spring 2019, Taught by Prof. Bin Yu. Helped in *redesigning* the class along with one other TA, Yuansi Chen, for a class of 140+ students.
- *Introduction to Machine Learning*: EECS 189, Spring 2018, Taught by Prof. Anant Sahai, and Prof. Jennifer Listgarten. *Co-led* the content development (homeworks, discussions, and exams) in a team of 20+ TAs in a class of 350+ students.

Teaching Assistant, IIT Bombay, and Government of India 2011–2014

Responsible for weekly discussions (40 students) besides exam grading, and (optional) help sessions for larger groups (up to 200 students)

- *Linear Algebra*: MA 106, Spring 2012, Spring 2013, Spring 2014 (for IIT Bombay, and separately for Ministry of Human Resource Development, Government of India's online course for 400 undergraduate colleges)
- *Ordinary Differential Equations*: MA 108, Spring 2012, Spring 2013
- *Partial Differential Equations*: MA 207, Summer 2012
- *Calculus*: MA 105, Fall 2011, Fall 2012 (Intensive Program for Entrants)
- *Electromagnetism*: PH 103, Fall 2011

Guest lectures

- *STAT 234, Harvard University*: Regret analysis of the posterior sampling (a series of 3 lectures) with titles: (i) The fellowship of Bayesian regret, (ii) the two Bellman operators, and (iii) the return of the martingale Apr 2022

ACADEMIC
SERVICES

- *STAT 234, Harvard University*: Offline off-policy reinforcement learning *Feb 2022*
- *STAT 212, UC Berkeley*: Revisiting Complexity, and the Bias-Variance Tradeoff: Using Minimum Description Length *Apr 2021*
- *EECS 189, UC Berkeley*: An Introduction to Ensemble Methods *Oct 2019*
- *STAT 154, UC Berkeley*: Boosting *Apr 2019*

Scientific Meetings

- Chair, New Researchers Group Session, IMS Annual Meeting *June 2022*
- Chair, Statistical Machine Learning Session, IMS Annual Meeting *June 2022*
- Mentor, Summer Institute on Just-in-Time Adaptive Interventions via MRTs *Oct 2021*

Institutional Mentoring Activities

- MIT Institute for Data, Systems, & Society (IDSS) Postdoc Mentors for *PhDs* *2022—*
- UC Berkeley Artificial Intelligence Research (BAIR) Buddies for *incoming PhDs* *2020—2021*
- UC Berkeley BAIR Undergraduate Mentoring Program for *undergraduates* *2017—2021*
- IIT Bombay Student Mentoring Program (ISMP) for *incoming undergraduates* *2013—2014*
- IIT Bombay, Academic Mentoring Program (DAMP) for *sophomores & juniors* *2012—2014*
- Intensive Program for Entrants (IPE) for *incoming undergraduates*, IIT Bombay *2012—2013*

Reviewing Activities

- *Journals*: Journal of Machine Learning Research (JMLR) (4 Papers), IEEE Transactions on Information Theory (4 Papers), Bernoulli (2 Papers), Mathematics of Operations Research (1 paper), Statistics, and Computing (1 paper), JRSSB (1 paper)
- *Conferences*: Conference on Learning Theory (COLT), International Conference on Machine Learning (ICML), Neural Information Processing Systems (NeurIPS), Foundations of Computer Science (FOCS), Symposium on Discrete Algorithms (SODA), AAAI Conference on Artificial Intelligence

Graduate Admissions

- EECS Graduate Admissions Committee, MIT *2021*
- EECS Graduate Admissions Committee, UC Berkeley *2018—2020*