

# Near-Optimal Compression In Near-Linear Time

**Raaz Dwivedi**, Lester Mackey, Abhishek Shetty  
[raaz@mit.edu](mailto:raaz@mit.edu)

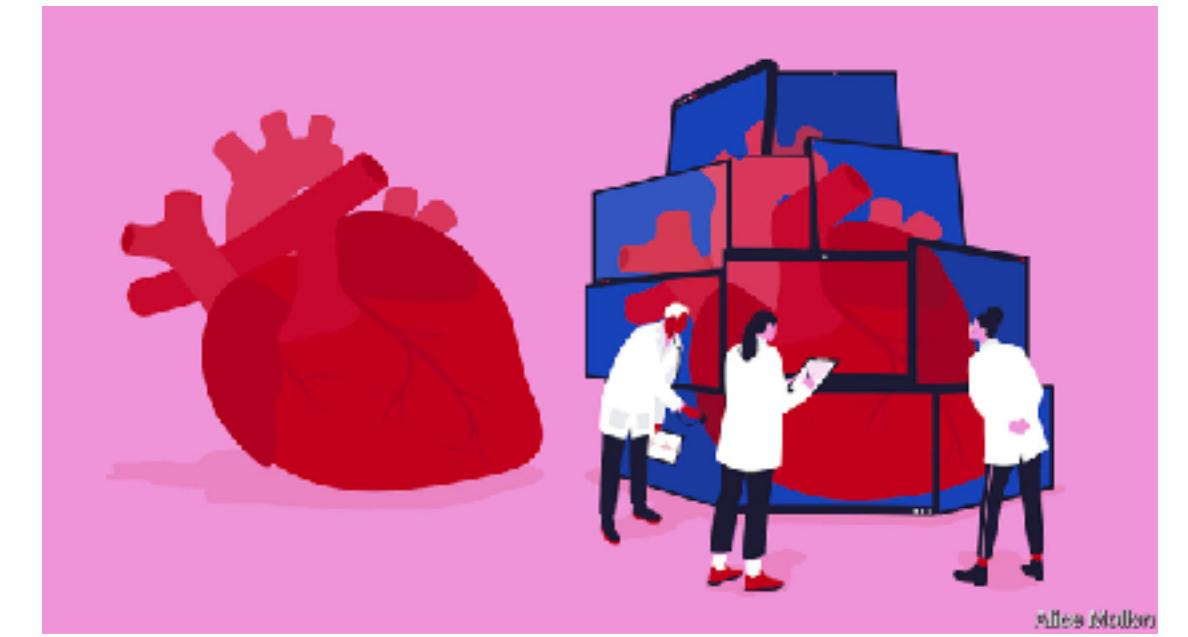


Harvard John A. Paulson  
School of Engineering  
and Applied Sciences



MIT LIDS Student Conference  
Jan 26, 2022

# Computational cardiology



- Modeling **digital twin heart** to predict therapy response in a *non-invasive* way requires single-cell modeling.

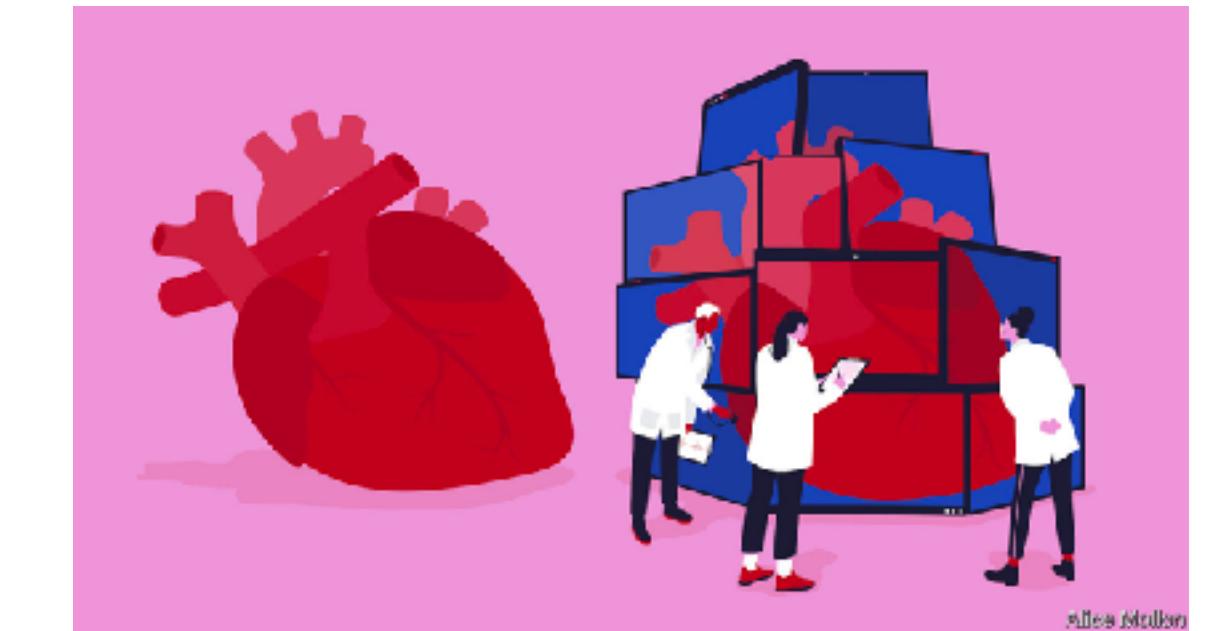
Commonly used strategy:

- Estimate single cell model using Bayesian set-up: Use millions of Markov chain Monte Carlo (MCMC) points to approximate posterior  $\mathbb{P}^*$
- Feed these samples to whole-heart simulator

$$\mathbb{P}^* f \triangleq \int f(x) d\mathbb{P}^*(x) \approx \frac{1}{n} \sum_{i=1}^n f(x_i) \triangleq \mathbb{P}_n f$$

$x_i$  = MCMC sample for single cell  
model parameters  
 $f$  = heart simulator

# Computational cardiology



- Modeling **digital twin heart** to predict therapy response in a *non-invasive way* requires single-cell modeling.

Commonly used strategy:

- Estimate single cell model using Bayesian set-up: Use millions of Markov chain Monte Carlo (MCMC) points to approximate posterior  $\mathbb{P}^*$
- Feed these samples to whole-heart simulator

$$\mathbb{P}^* f \triangleq \int f(x) d\mathbb{P}^*(x) \approx \frac{1}{n} \sum_{i=1}^n f(x_i) \triangleq \mathbb{P}_n f$$

$x_i$  = MCMC sample for single cell model parameters  
 $f$  = heart simulator

1 Million MCMC samples ~ 2 weeks  
Single evaluation of  $f$  ~ 5 weeks

Can NOT use all million samples....!

# Goal: Represent $\mathbb{P}^*$ using a few high quality points $(x_i)_{i=1}^n$

- Typical solutions like **I.I.D. sampling, and MCMC sampling** exhibit the **slow root-n Monte Carlo rate**  $|\mathbb{P}^*f - \mathbb{P}_n f| = \Theta(n^{-1/2})$ , e.g., around a million points for 0.1% error
- Computationally prohibitive for expensive  $f$ , and common fixes uniform thinning, or **standard thinning--choose every  $t$ -th point**
- **Accuracy degrades with such thinning**— $\Theta(\sqrt{t/n})$  worst-case error—same as the Monte Carlo rate with  $n/t$  points;  
e.g.,  $n^{-1/4}$  rate with  $\sqrt{n}$  points

Can we do better?

# Minimax lower bounds

- There exists some  $\mathbb{P}^*$  such that the worst-case integration error
  - Is  $\Omega(n^{-1/2})$  for any compression scheme returning  $\sqrt{n}$  points [Philips and Tai, 2020]

# Minimax lower bounds

- There exists some  $\mathbb{P}^*$  such that the worst-case integration error
  - Is  $\Omega(n^{-1/2})$  for any compression scheme returning  $\sqrt{n}$  points  
[Philips and Tai, 2020]
  - Is  $\Omega(n^{-1/2})$  for any approximation based on  $\sqrt{n}$  i.i.d. points  
[Tolstikhin, Sriperumbudur, and Muandet, 2017]

# This talk: Kernel thinning-Compress++

- **KT-Compress++**: A practical strategy based on two new algorithms to provide near-optimal compression in near-linear time
  - **Kernel thinning (KT)**: Provides near-optimal compression
  - **Compress++**: Provides significantly reduced runtime for generic thinning algorithms with minimal worsening of error

# This talk: Kernel thinning-Compress++

- **KT-Compress++**: A practical strategy based on two new algorithms to provide near-optimal compression in near-linear time
  - **Kernel thinning (KT)**: Provides near-optimal compression
  - **Compress++**: Provides significantly reduced runtime for generic thinning algorithms with minimal worsening of error
- Overall, KT-Compress++ a solution for finding
  - better than Monte Carlo points
  - high quality coresets
  - good prototypes

# Problem set-up

- **Input:**

Points  $(x_i)_{i=1}^n \subset \mathbb{R}^d$  with empirical distribution  $\mathbb{P}_{in} \triangleq \frac{1}{n} \sum_{i=1}^n \delta_{x_i}$

Target output size  $s$  ( $s = \sqrt{n}$  for heavy compression)

- **Goal:**

Return a subset of input points with size  $s$ , empirical distribution  $\mathbb{P}_{out}$  with error rate  $o(s^{-1/2})$ , i.e., better than Monte Carlo rate

# Reproducing kernel Hilbert space (RKHS)

- RKHS of  $\mathbf{k}$  is given by  $\mathbb{H}_{\mathbf{k}} \triangleq \overline{\text{span}}\{\mathbf{k}(x, \cdot), x \in \mathbb{R}^d\}$
- $\mathbf{k}$  is universal when  $\mathbb{H}_{\mathbf{k}}$  is dense in the space of continuous functions, e.g.,

$$\text{Gaussian } \mathbf{k}(x, y) = \exp\left(-\frac{1}{2}\|x - y\|^2\right); \text{ IMQ } \mathbf{k}(x, y) = \frac{1}{(1 + \|x - y\|_2^2)^{1/2}}$$

$\mathbf{k} : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  is a reproducing kernel if the matrix  $\mathbf{K} = (\mathbf{k}(x_i, x_j))_{i,j=1}^n$  is a symmetric positive definite matrix for any  $n$  and any  $(x_1, \dots, x_n)$

# Kernel Thinning

$x_1, x_2, \dots, x_n \in \mathbb{R}^d$   
smooth decaying  $\mathbf{k}$

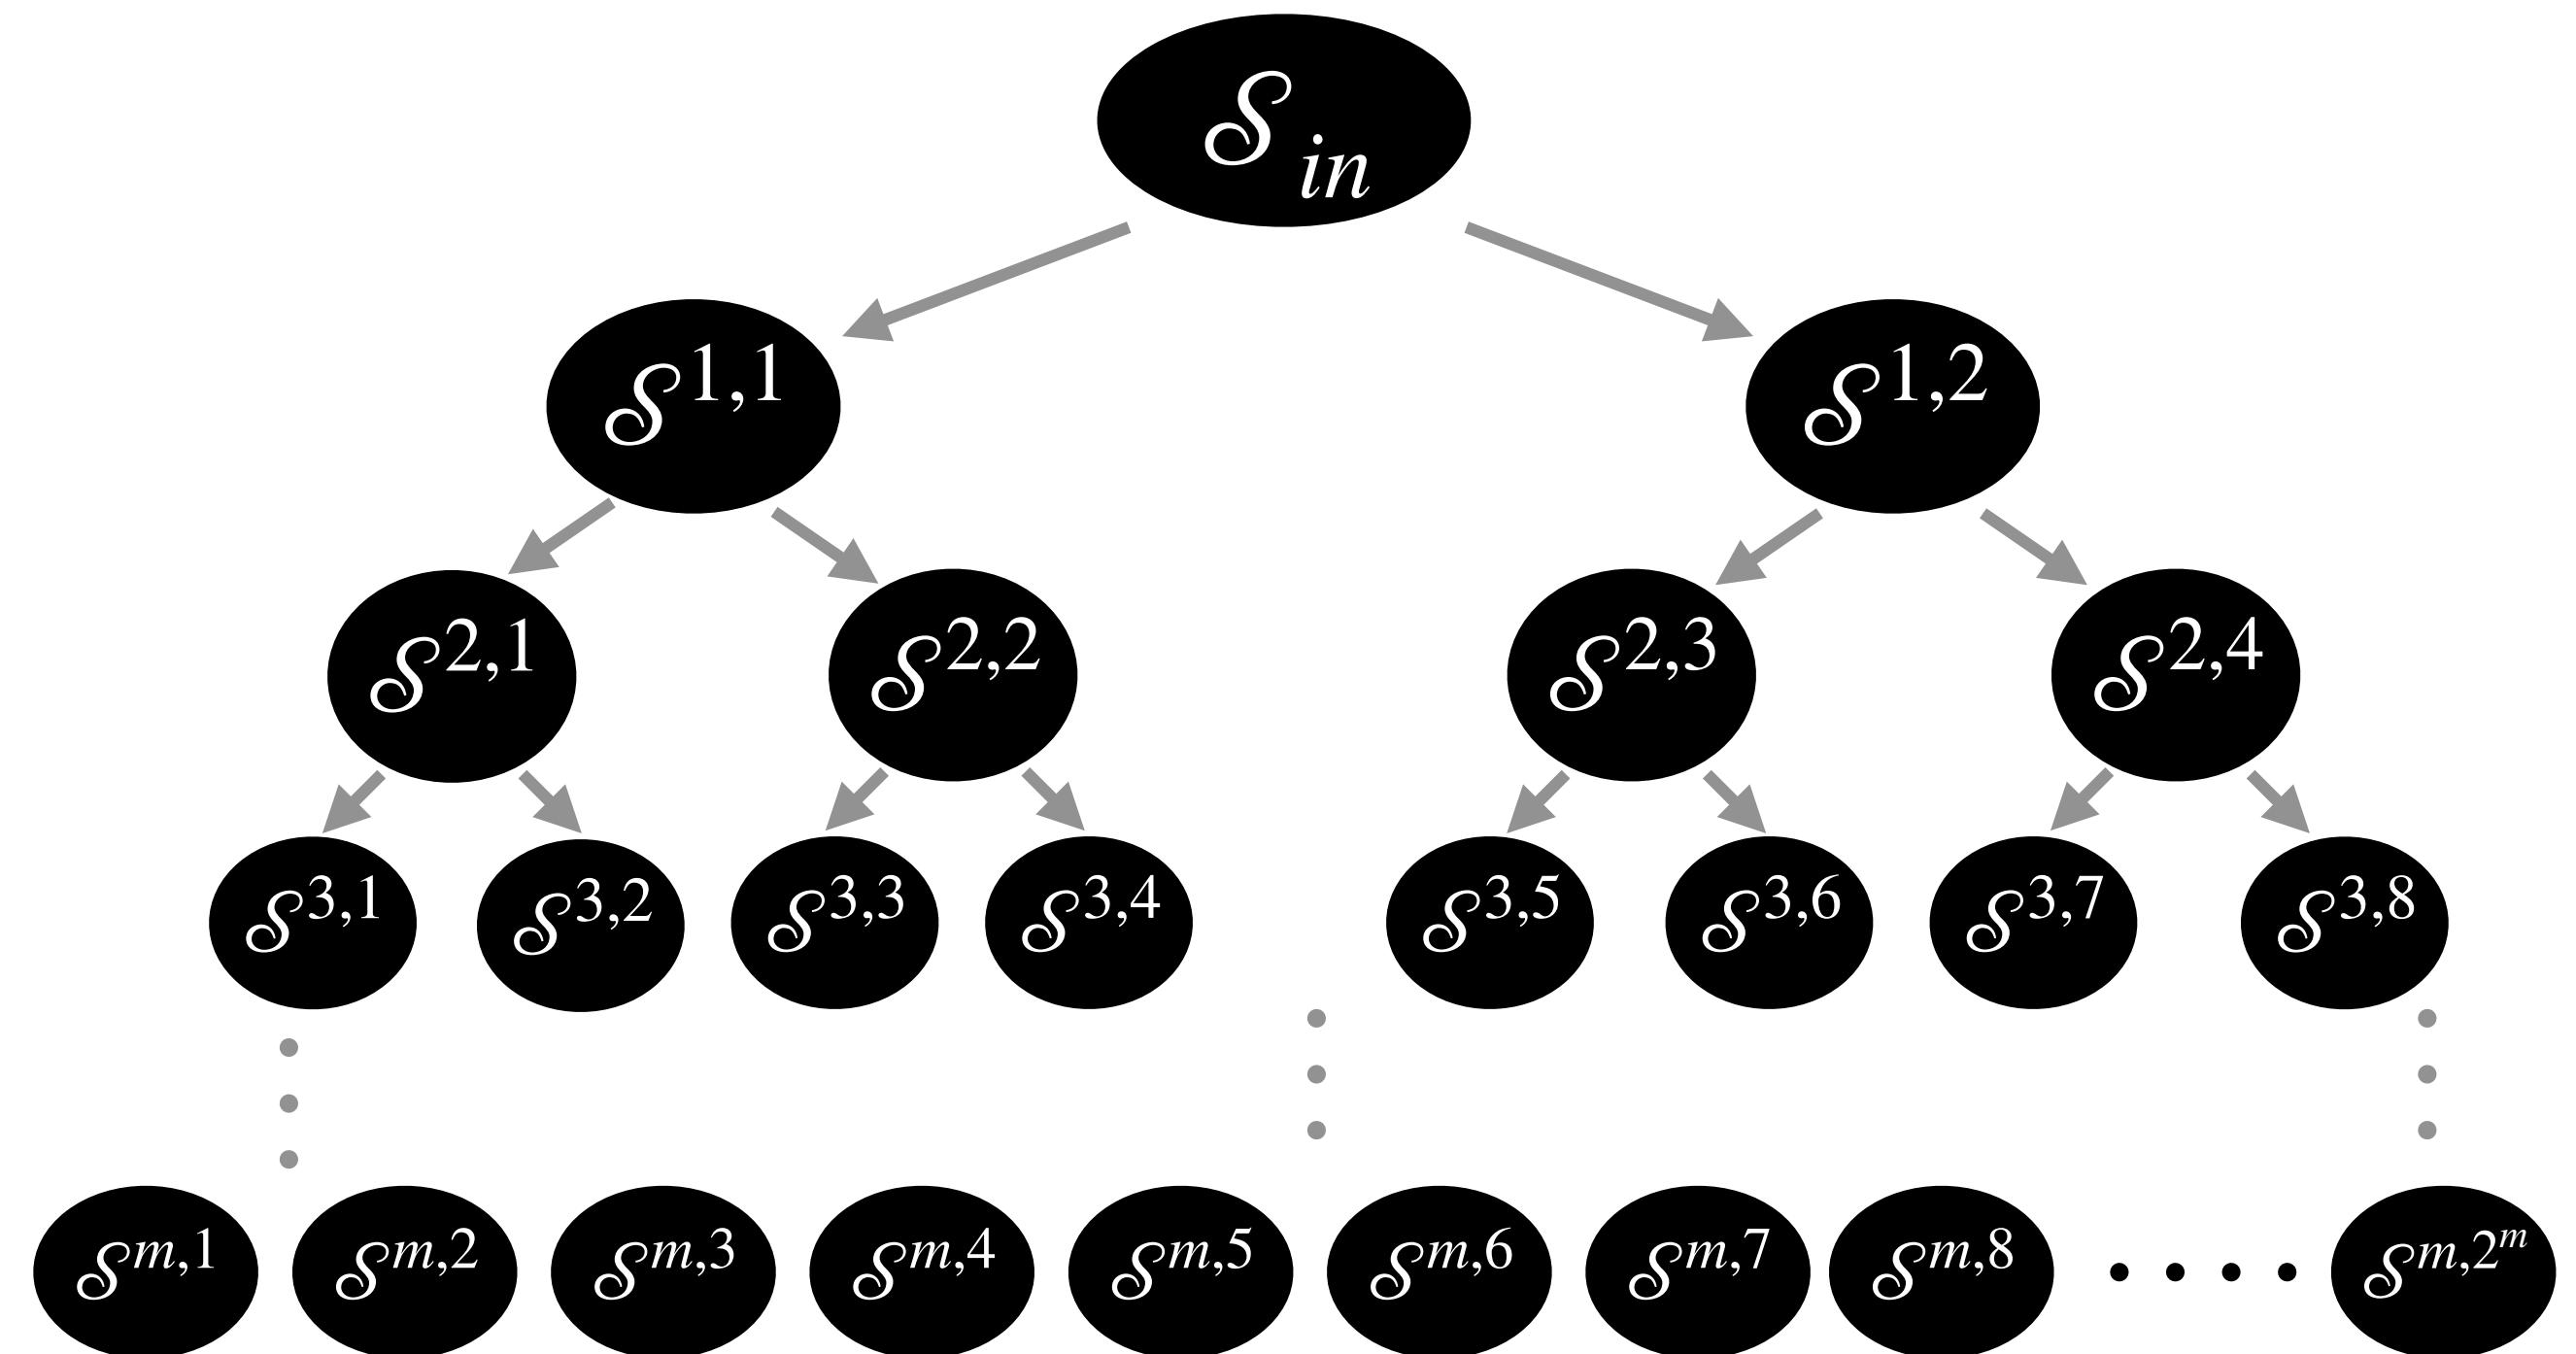
$$\mathbb{P}_{in} \triangleq \frac{1}{n} \sum_{i=1}^n \delta_{x_i}$$



Non-uniform sub-sample of **size  $s$**

$$y_1, \dots, y_s \\ \mathbb{P}_{KT} \triangleq \frac{1}{s} \sum_{i=1}^s \delta_{y_i}$$

- **Stage 1:** Repeated rounds of **non-uniform splitting** the parent coresset in two equal-sized children coressets
- **Stage 2:** Point-by-point **refinement** of the best child coreset



# Kernel Thinning: Better than Monte Carlo rate

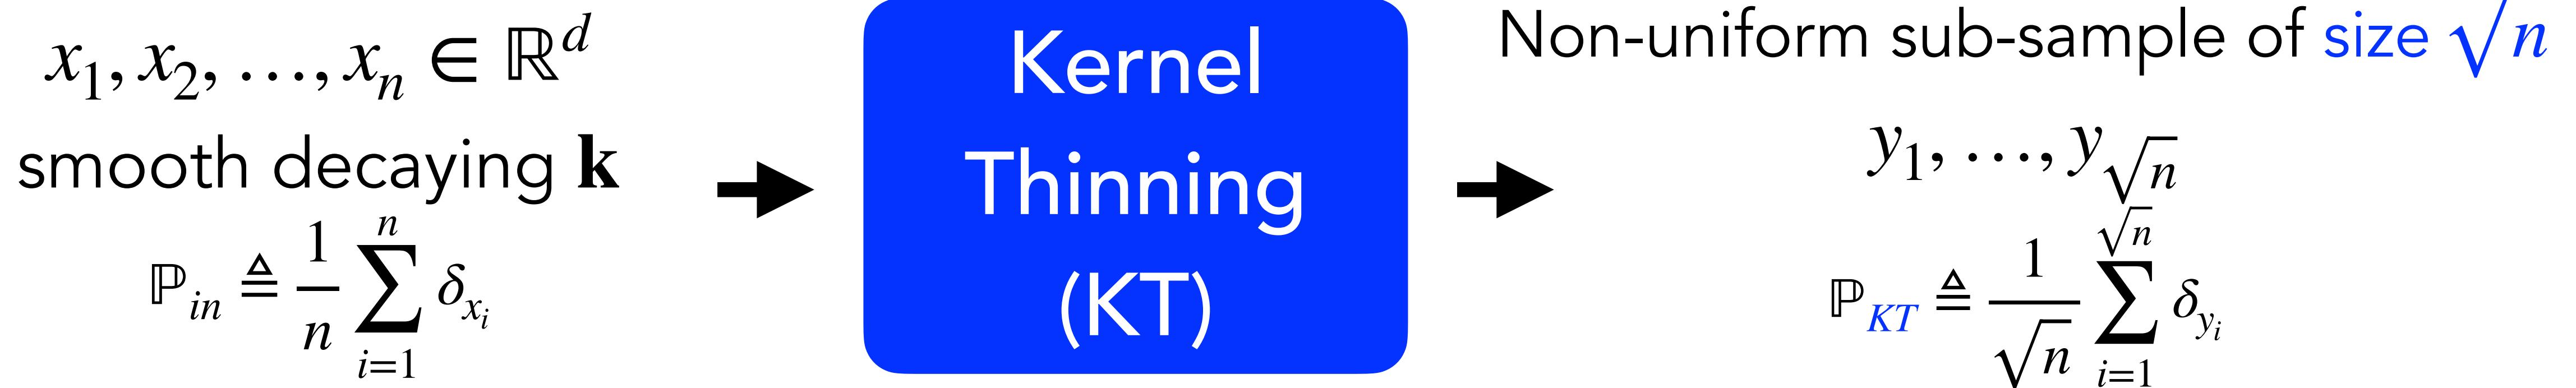


For any fixed  $g \in \mathbb{H}_k$ , with high probability over the randomness in KT, we have

$$|\mathbb{P}_{in}g - \mathbb{P}_{KT}g| \lesssim \frac{1}{s} \cdot \|g\|_k \sqrt{\|k\|_\infty (\log s + \log \log(n/s))}$$

Much faster than the Monte Carlo rate for standard/uniform thinning  $\mathcal{O}\left(\frac{1}{\sqrt{s}}\right)$

# Kernel Thinning: Better than Monte Carlo rate

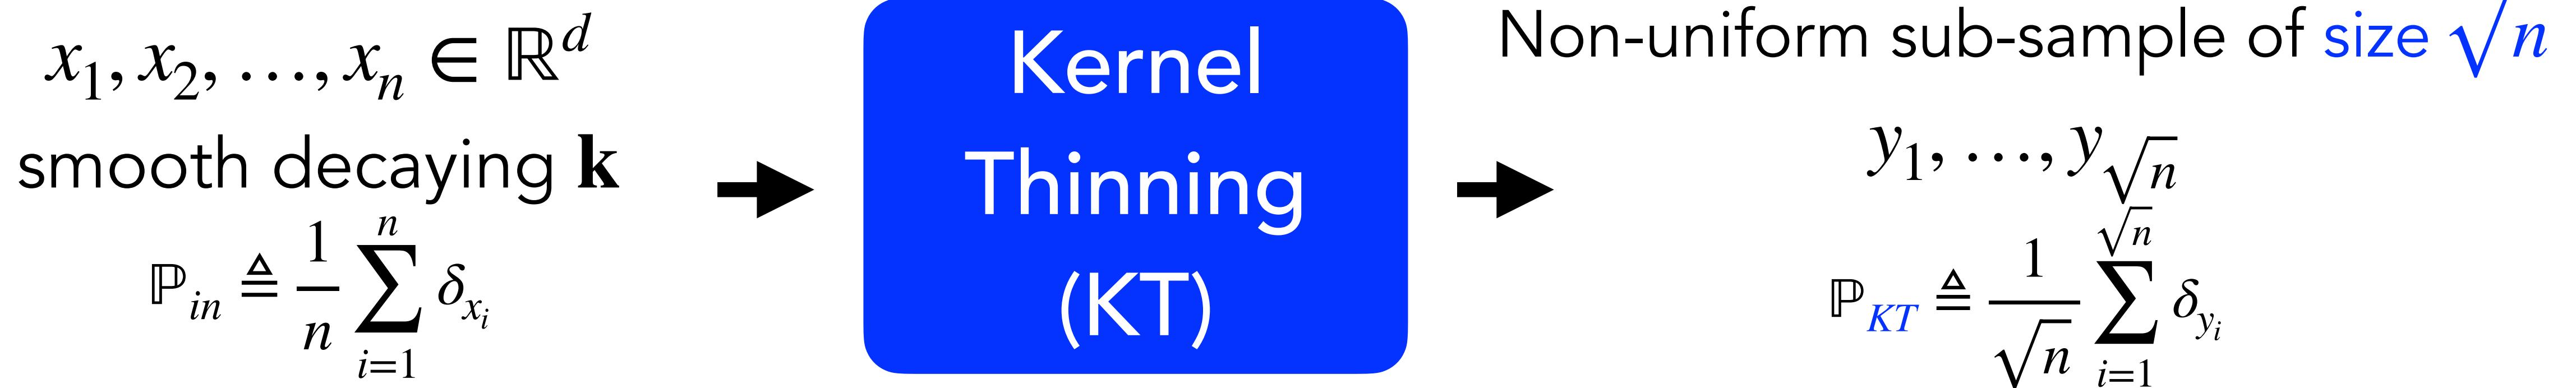


For any fixed  $g \in \mathbb{H}_k$ , with high probability over the randomness in KT, we have

$$|\mathbb{P}_{in}g - \mathbb{P}_{KT}g| \lesssim \frac{1}{\sqrt{n}} \cdot \|g\|_k \sqrt{\|\mathbf{k}\|_\infty \log n}$$

Much faster than the Monte Carlo rate for standard/uniform thinning  $\mathcal{O}\left(\frac{1}{n^{1/4}}\right)$

# Kernel Thinning: Better than Monte Carlo rate



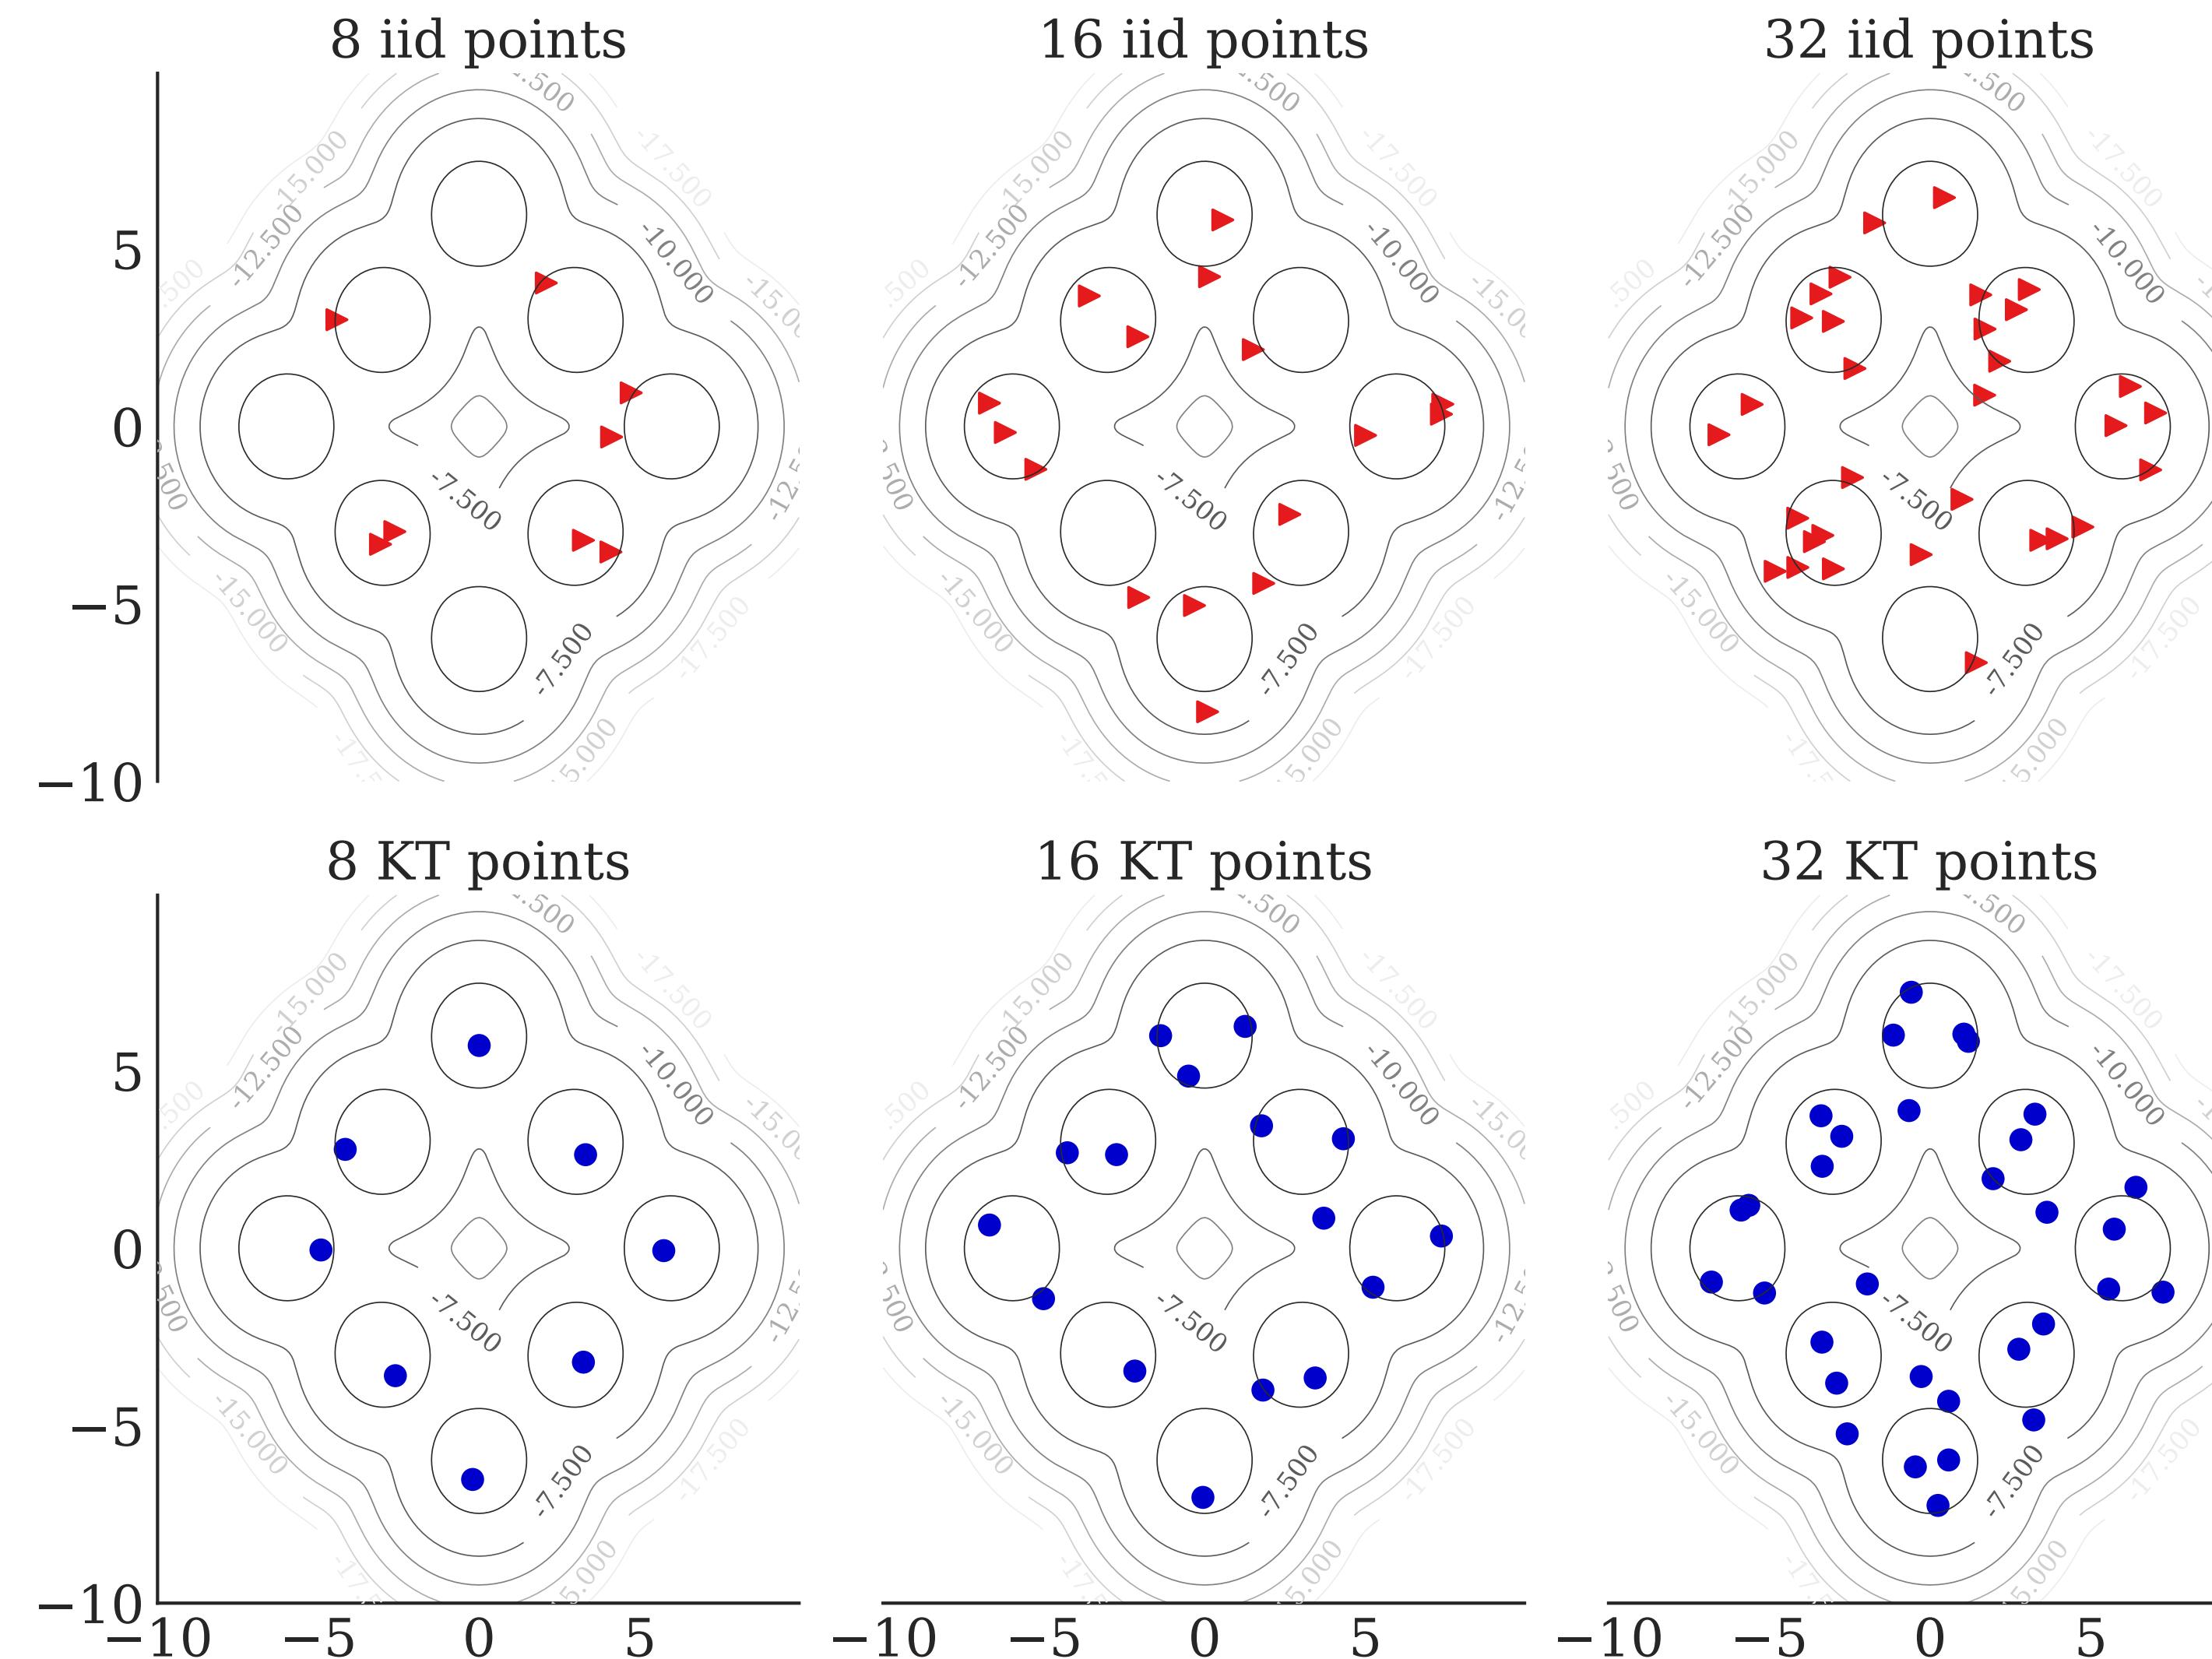
For any fixed  $g \in \mathbb{H}_k$ , with high probability over the randomness in KT, we have

$$|\mathbb{P}_{in}g - \mathbb{P}_{KT}g| \lesssim \frac{1}{\sqrt{n}} \cdot \|g\|_k \sqrt{\|\mathbf{k}\|_\infty \log n}$$

$$|\mathbb{P}^\star g - \mathbb{P}_{KT}g| = \mathcal{O}\left(\sqrt{\frac{\log n}{n}}\right) \text{ whenever } |\mathbb{P}^\star g - \mathbb{P}_{in}g| = \mathcal{O}\left(\sqrt{\frac{\log n}{n}}\right)$$

Easily satisfied for  
input from iid  
sampling, MCMC,  
quadrature  
methods....

# Kernel Thinning: Finds “diverse” set of points



# Worst-case error: $\widetilde{O}(n^{-1/2})$ with $\sqrt{n}$ points

- With high probability, the worst-case error--MMD error--in the reproducing kernel Hilbert space (RKHS) satisfies

$$\sup_{\|g\|_k \leq 1} |\mathbb{P}_{in}g - \mathbb{P}_{KT}g| \lesssim_d \begin{cases} n^{-1/2}\sqrt{\log n} & \text{(Compactly supported; e.g., B-spline } k) \\ n^{-1/2}\sqrt{\log^{d/2+1} n \log \log n} & \text{(Sub-Gaussian tails; e.g., Gaussian } k) \\ n^{-1/2}\sqrt{\log^{d+1} n \log \log n} & \text{(Sub-exponential tails; e.g., Matern } k) \end{cases}$$

- For output size  $s$ , the MMD error is  $\widetilde{O}(1/s)$

# Worst-case error: $\widetilde{O}(n^{-1/2})$ with $\sqrt{n}$ points

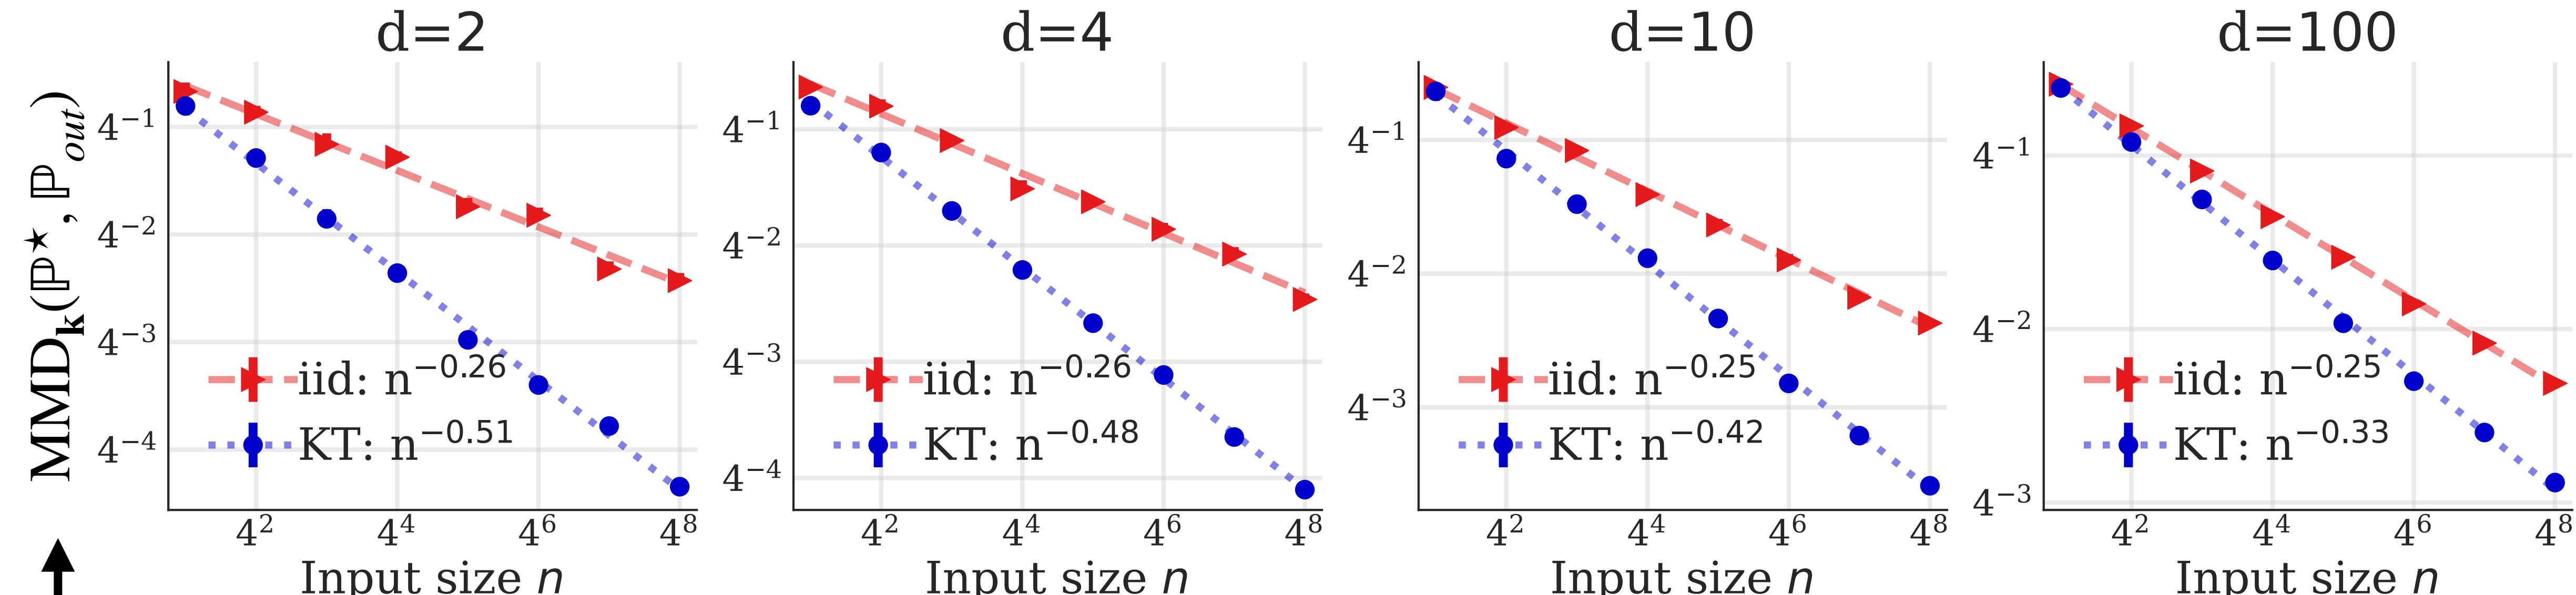
- With high probability, the worst-case error--MMD error--in the reproducing kernel Hilbert space (RKHS) satisfies

$$\sup_{\|g\|_k \leq 1} |\mathbb{P}^{\star} g - \mathbb{P}_{KT} g| \lesssim_d \begin{cases} n^{-1/2} \sqrt{\log n} & \text{(Compactly supported; e.g., B-spline } k) \\ n^{-1/2} \sqrt{\log^{d/2+1} n \log \log n} & \text{(Sub-Gaussian tails; e.g., Gaussian } k) \\ n^{-1/2} \sqrt{\log^{d+1} n \log \log n} & \text{(Sub-exponential tails; e.g., Matern } k) \end{cases}$$

whenever  $\sup_{\|g\|_k \leq 1} |\mathbb{P}^{\star} g - \mathbb{P}_{in} g| \lesssim n^{-1/2}$  --holds for iid / fast mixing MCMC input

# KT vs iid: Gaussian $P^*$ in $\mathbb{R}^d$

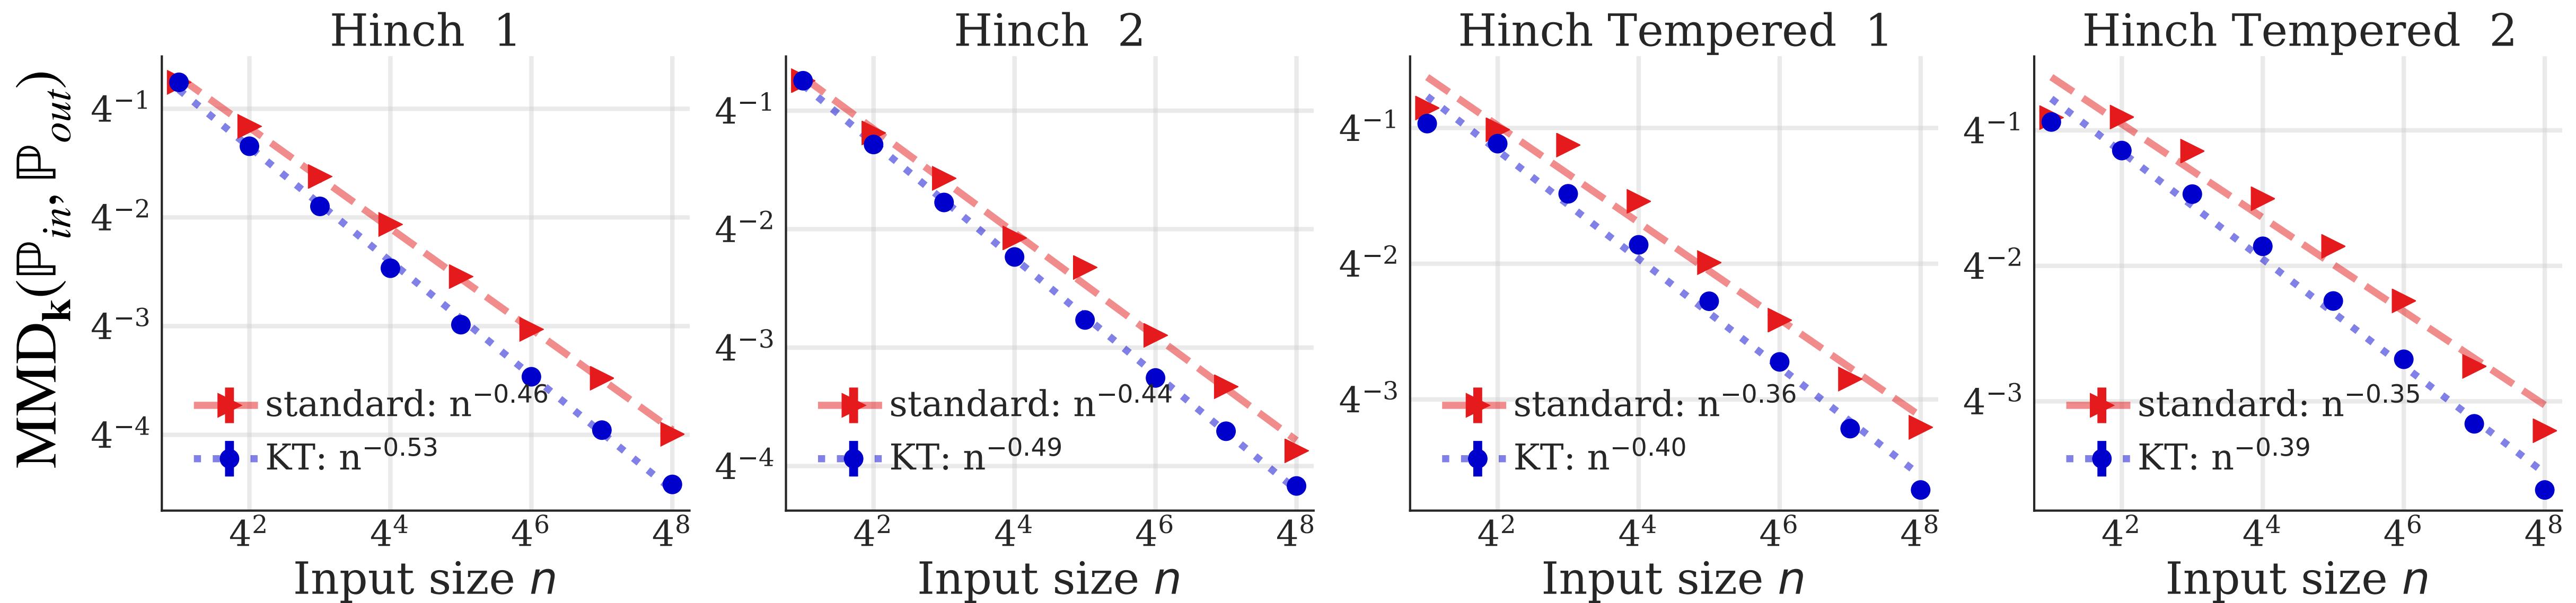
(Gaussian kernel with  $\sigma^2 = 2d$ )



In practice, significant gains even in dimension  $d = 100$

(Worst-case error in the unit ball of Gaussian RKHS)

# KT on MCMC samples from computational cardiology



In this setting with  $d = 38$ , standard thinning is already good,  
but KT provides further improvement!  
Note that each point saves 1000s of CPU hours!!

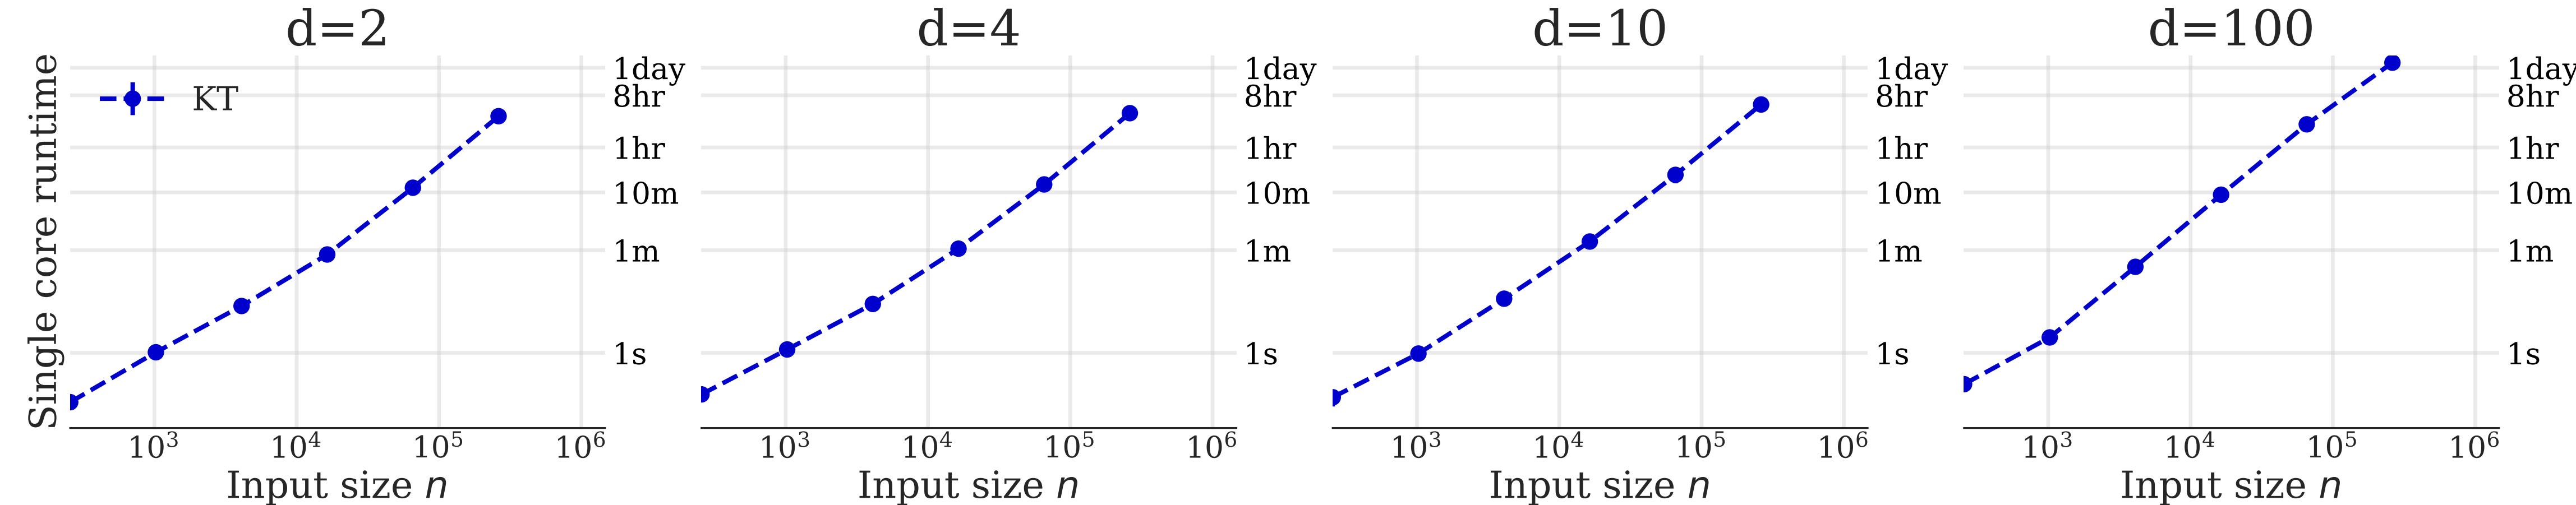
\*MCMC samples taken from Riabiz-Chen-Cockayne-Swietach-Niederer-Mackey-Oates, 2021

# Comparison with other works

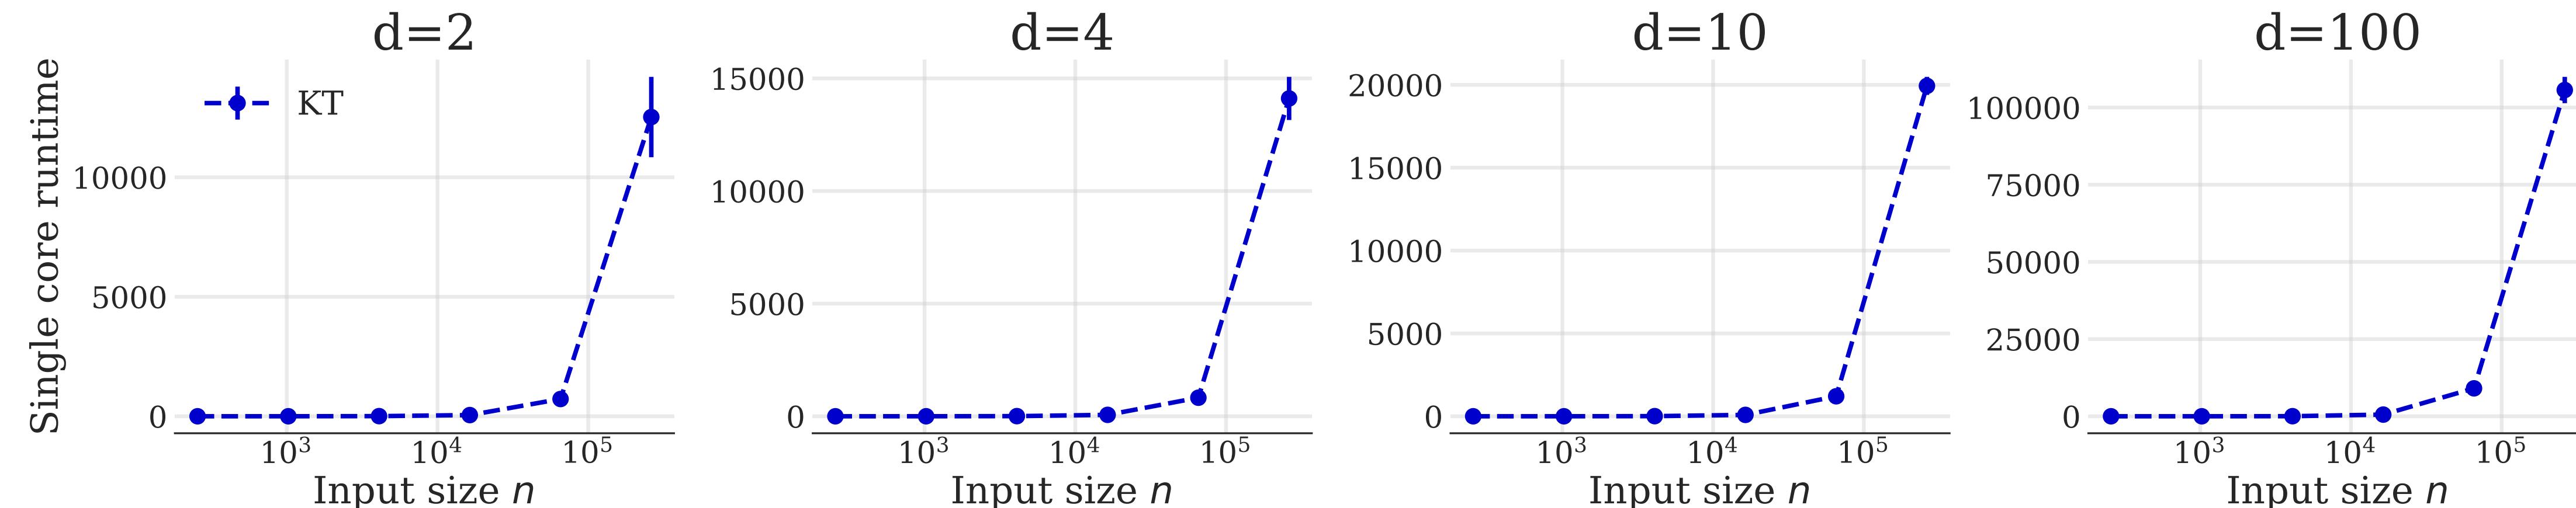
- Ours is the **first result** that establishes **faster than Monte Carlo rate for generic  $k$  and  $\mathbb{P}^*$  with unbounded support**
- Prior results show better than Monte Carlo rates
  - either empirically without a proof
  - or provide a proof under restrictive assumptions on  $\mathbb{P}^*$  or  $k$ , e.g., uniform distribution on unit cube, or finite-dimensional  $k$
  - including Kernel herding, quasi Monte Carlo, Bayesian quadrature, support points, stein points, ...

# KT drawback: $n^2$ runtime with $n$ input points

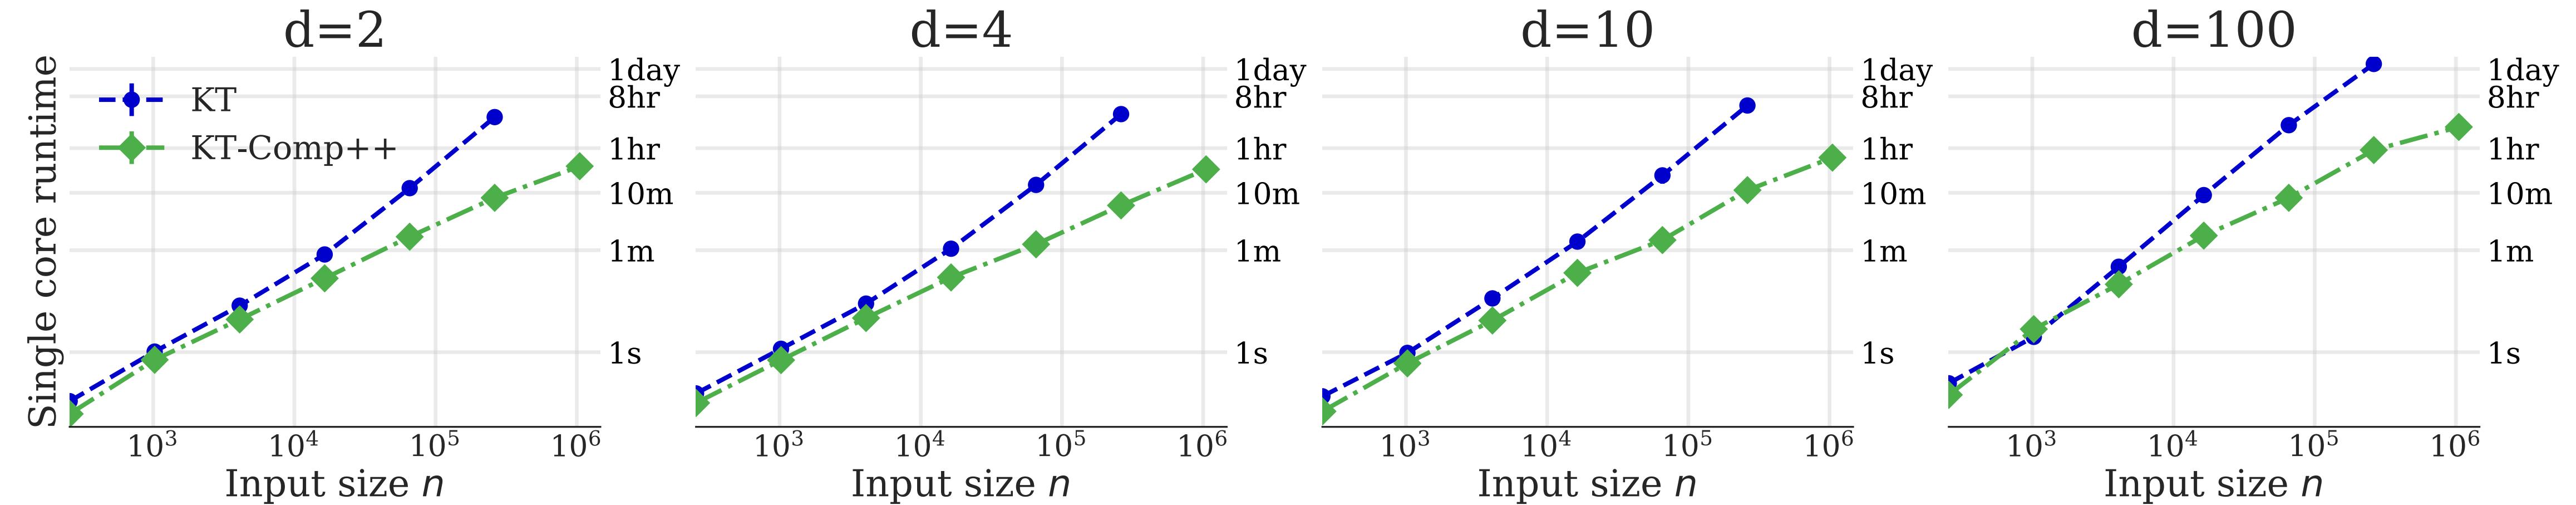
Y-axis = Runtime in log-scale



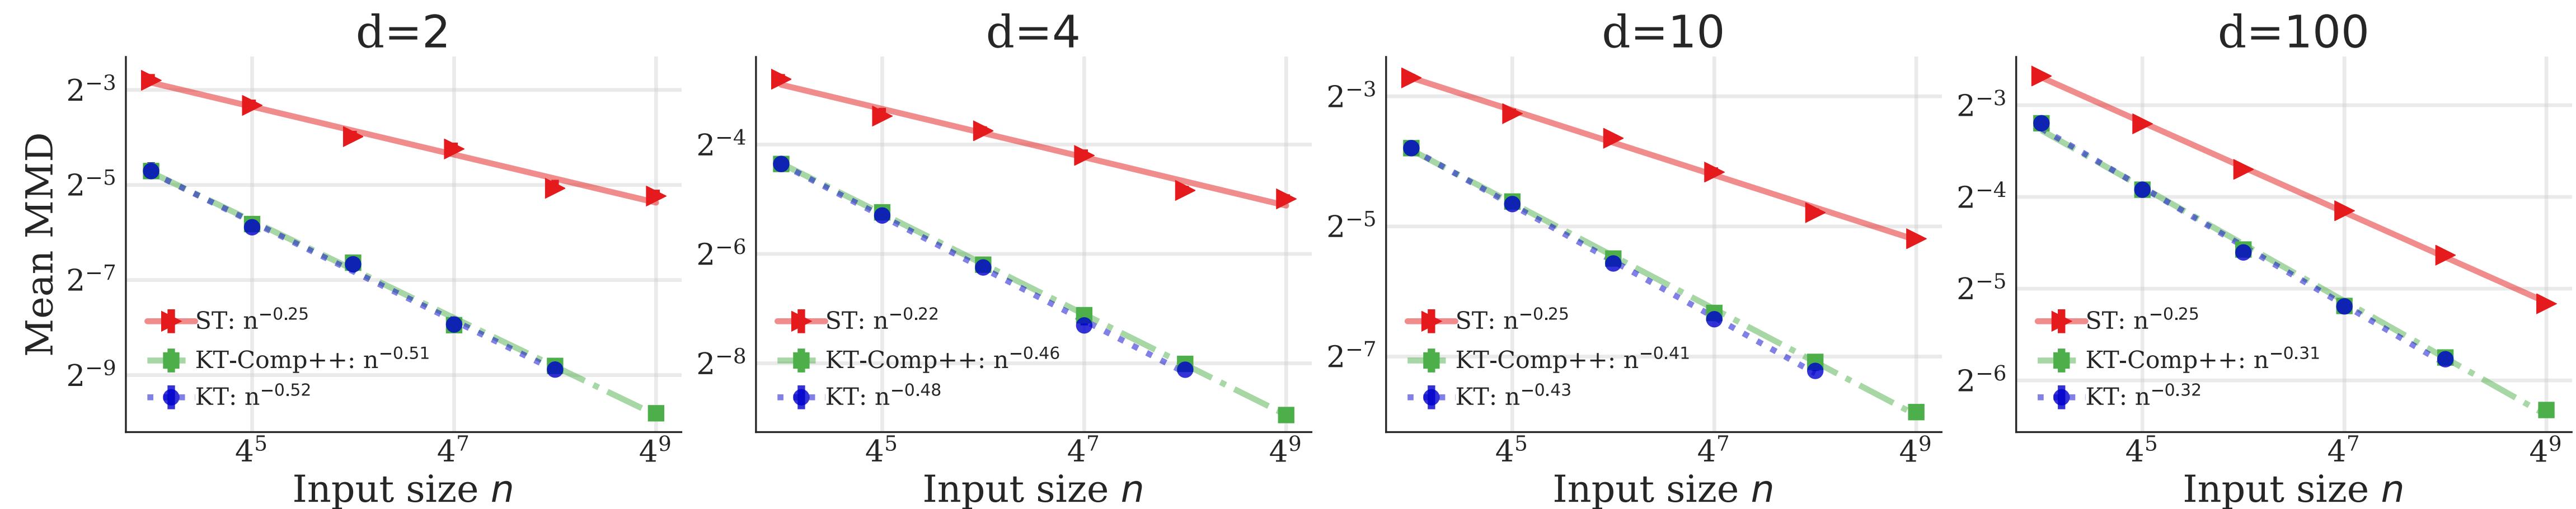
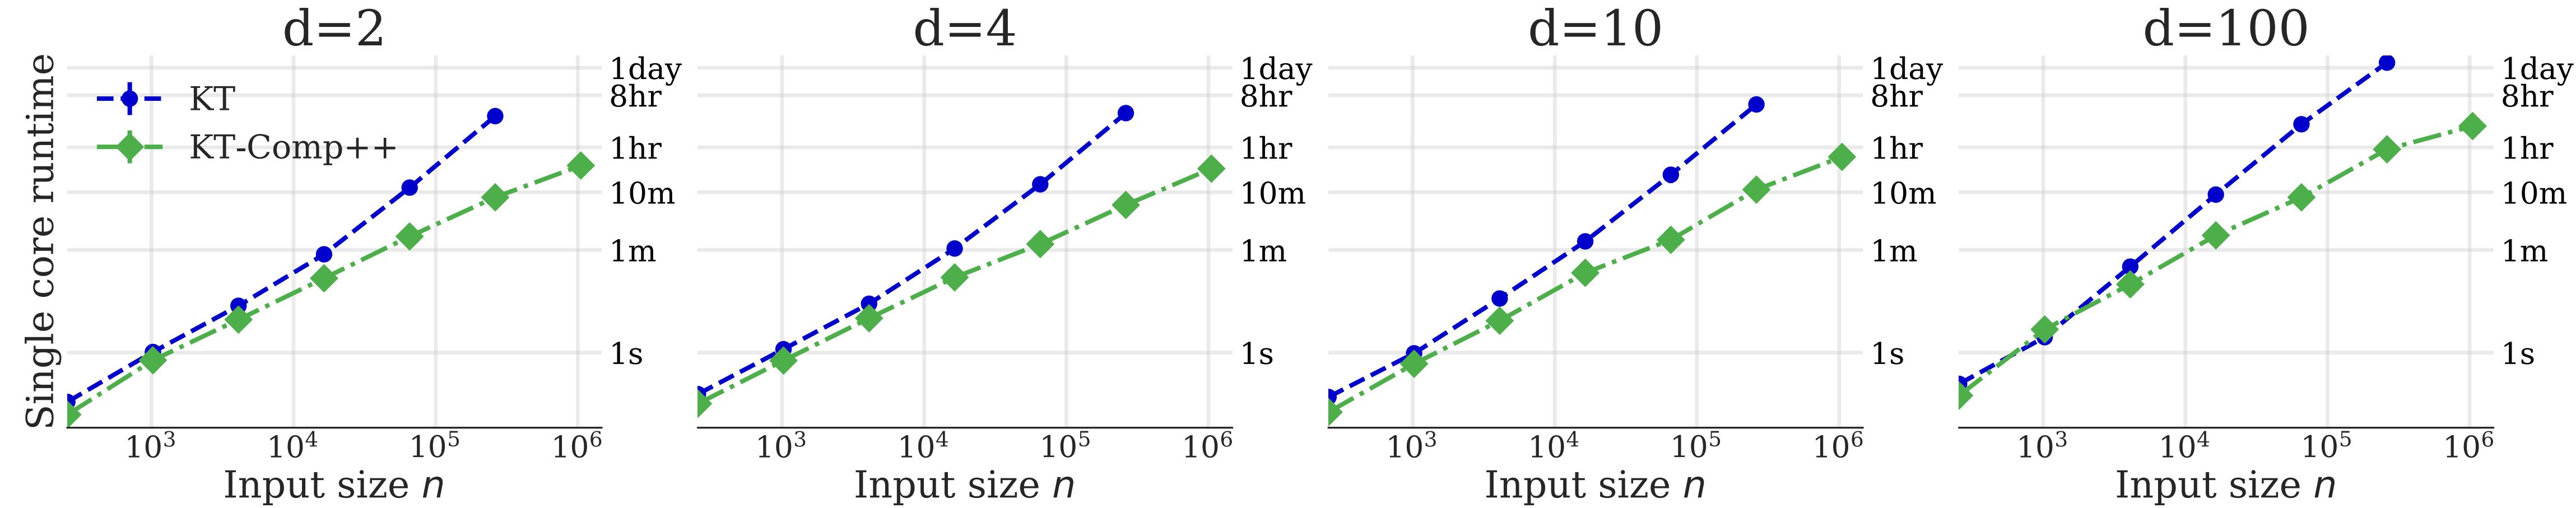
Y-axis = Runtime in linear scale (seconds)



# Compress++: Reducing runtime



# Compress++: Reducing runtime with minimal loss in accuracy!!



# Compress++: A recursive strategy to reduce runtime for generic thinning algorithms

Thinning algorithm with  
 $n^2$  runtime and  
error  $e(n)$

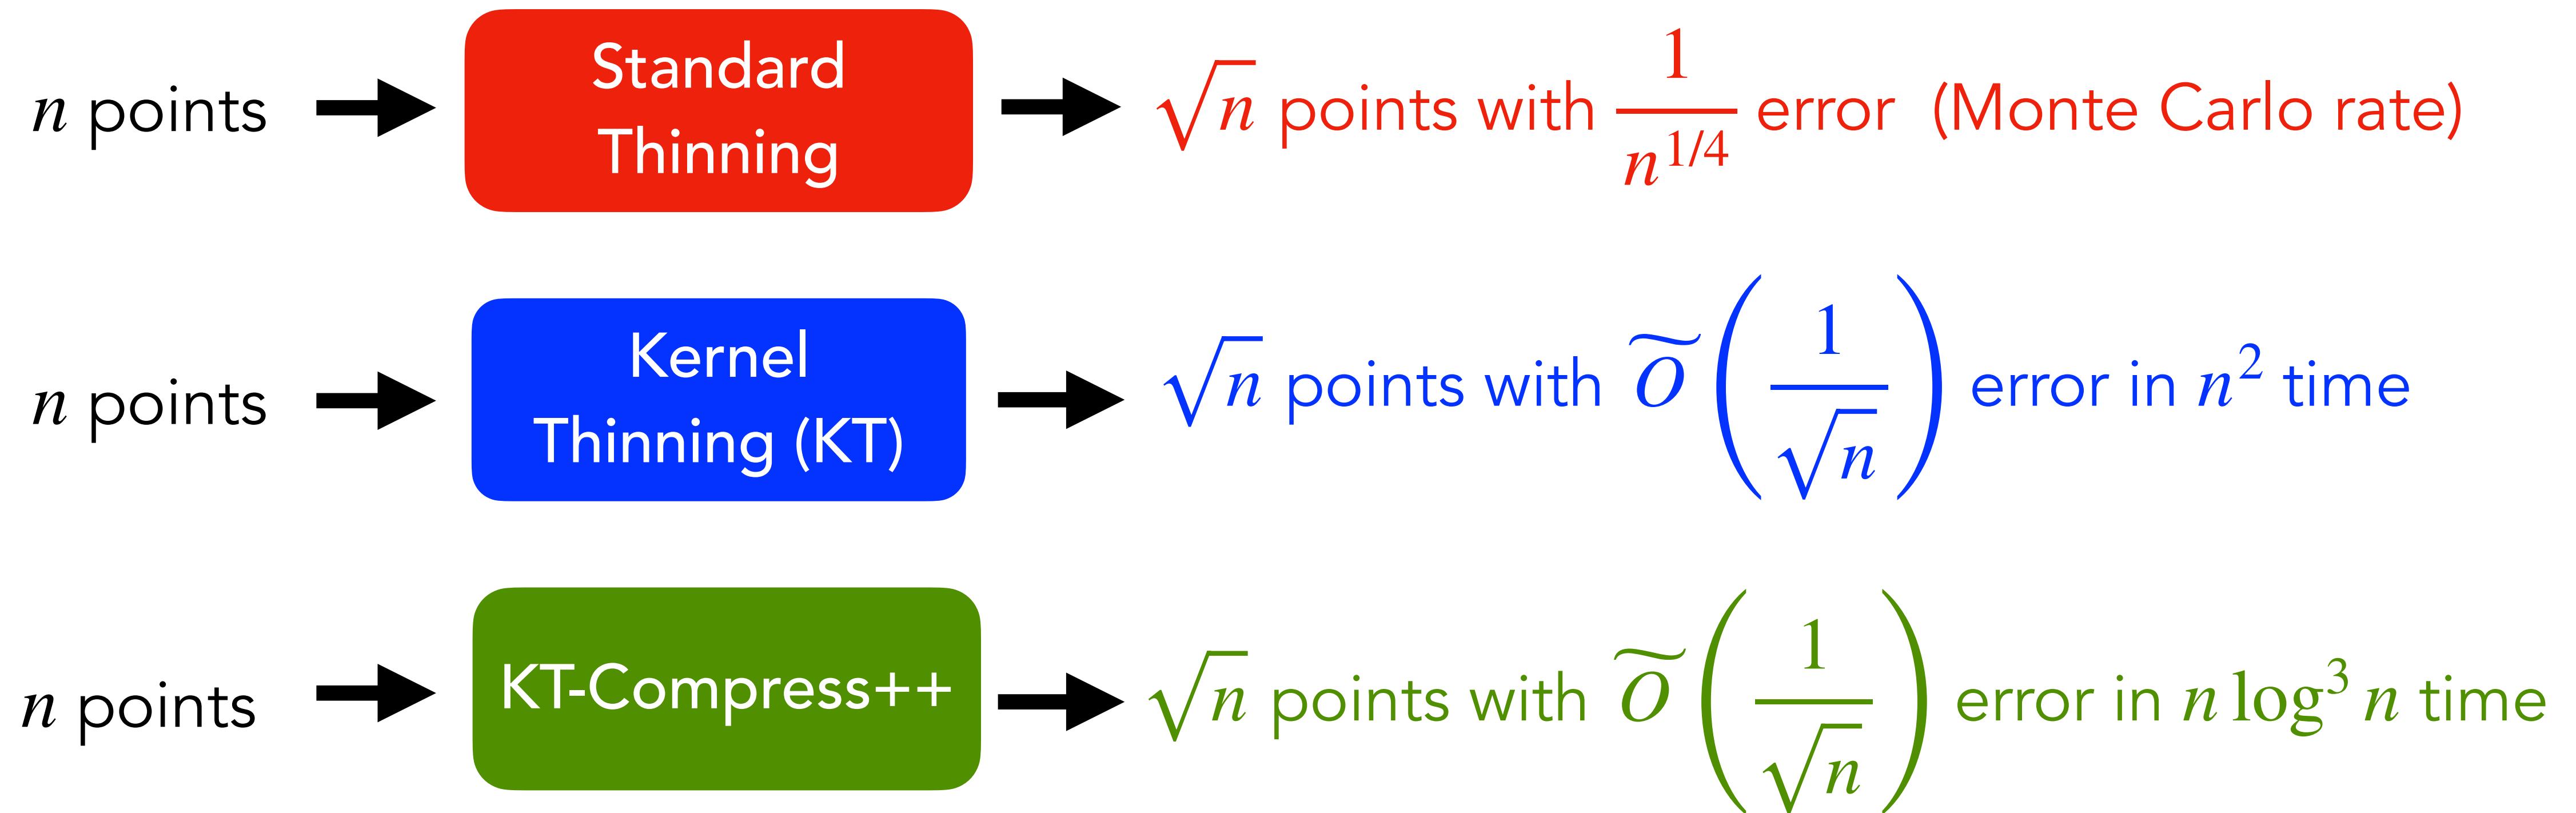


$\sqrt{n}$ -thinning algorithm with  
 $n \log^3 n$  runtime and  
error  $4e(n)$

- Compress( $\mathcal{S}$ , ALG):
  - If  $\text{size}(\mathcal{S}) == 1$ , **Return**  $\mathcal{S}$
  - Else:
    - (i) Call Compress separately on 4 equal splits of  $\mathcal{S}$
    - (ii) Concatenate the 4 outputs from step (i)
    - (iii) **Return** Halved output of step (ii) using ALG

## Summary:

KT-Compress++ provides near-optimal compression in near-linear time



 python™ pip install goodpoints

<https://arxiv.org/abs/2105.05842>

[Kernel Thinning, COLT 2021]

<https://arxiv.org/abs/2110.01593>

[Generalized Kernel Thinning, ICLR 2022]

<https://arxiv.org/pdf/2111.07941.pdf>

[Distribution Compression In Near-linear Time, ICLR 2022]

# Additional slides

# Details of kernel thinning

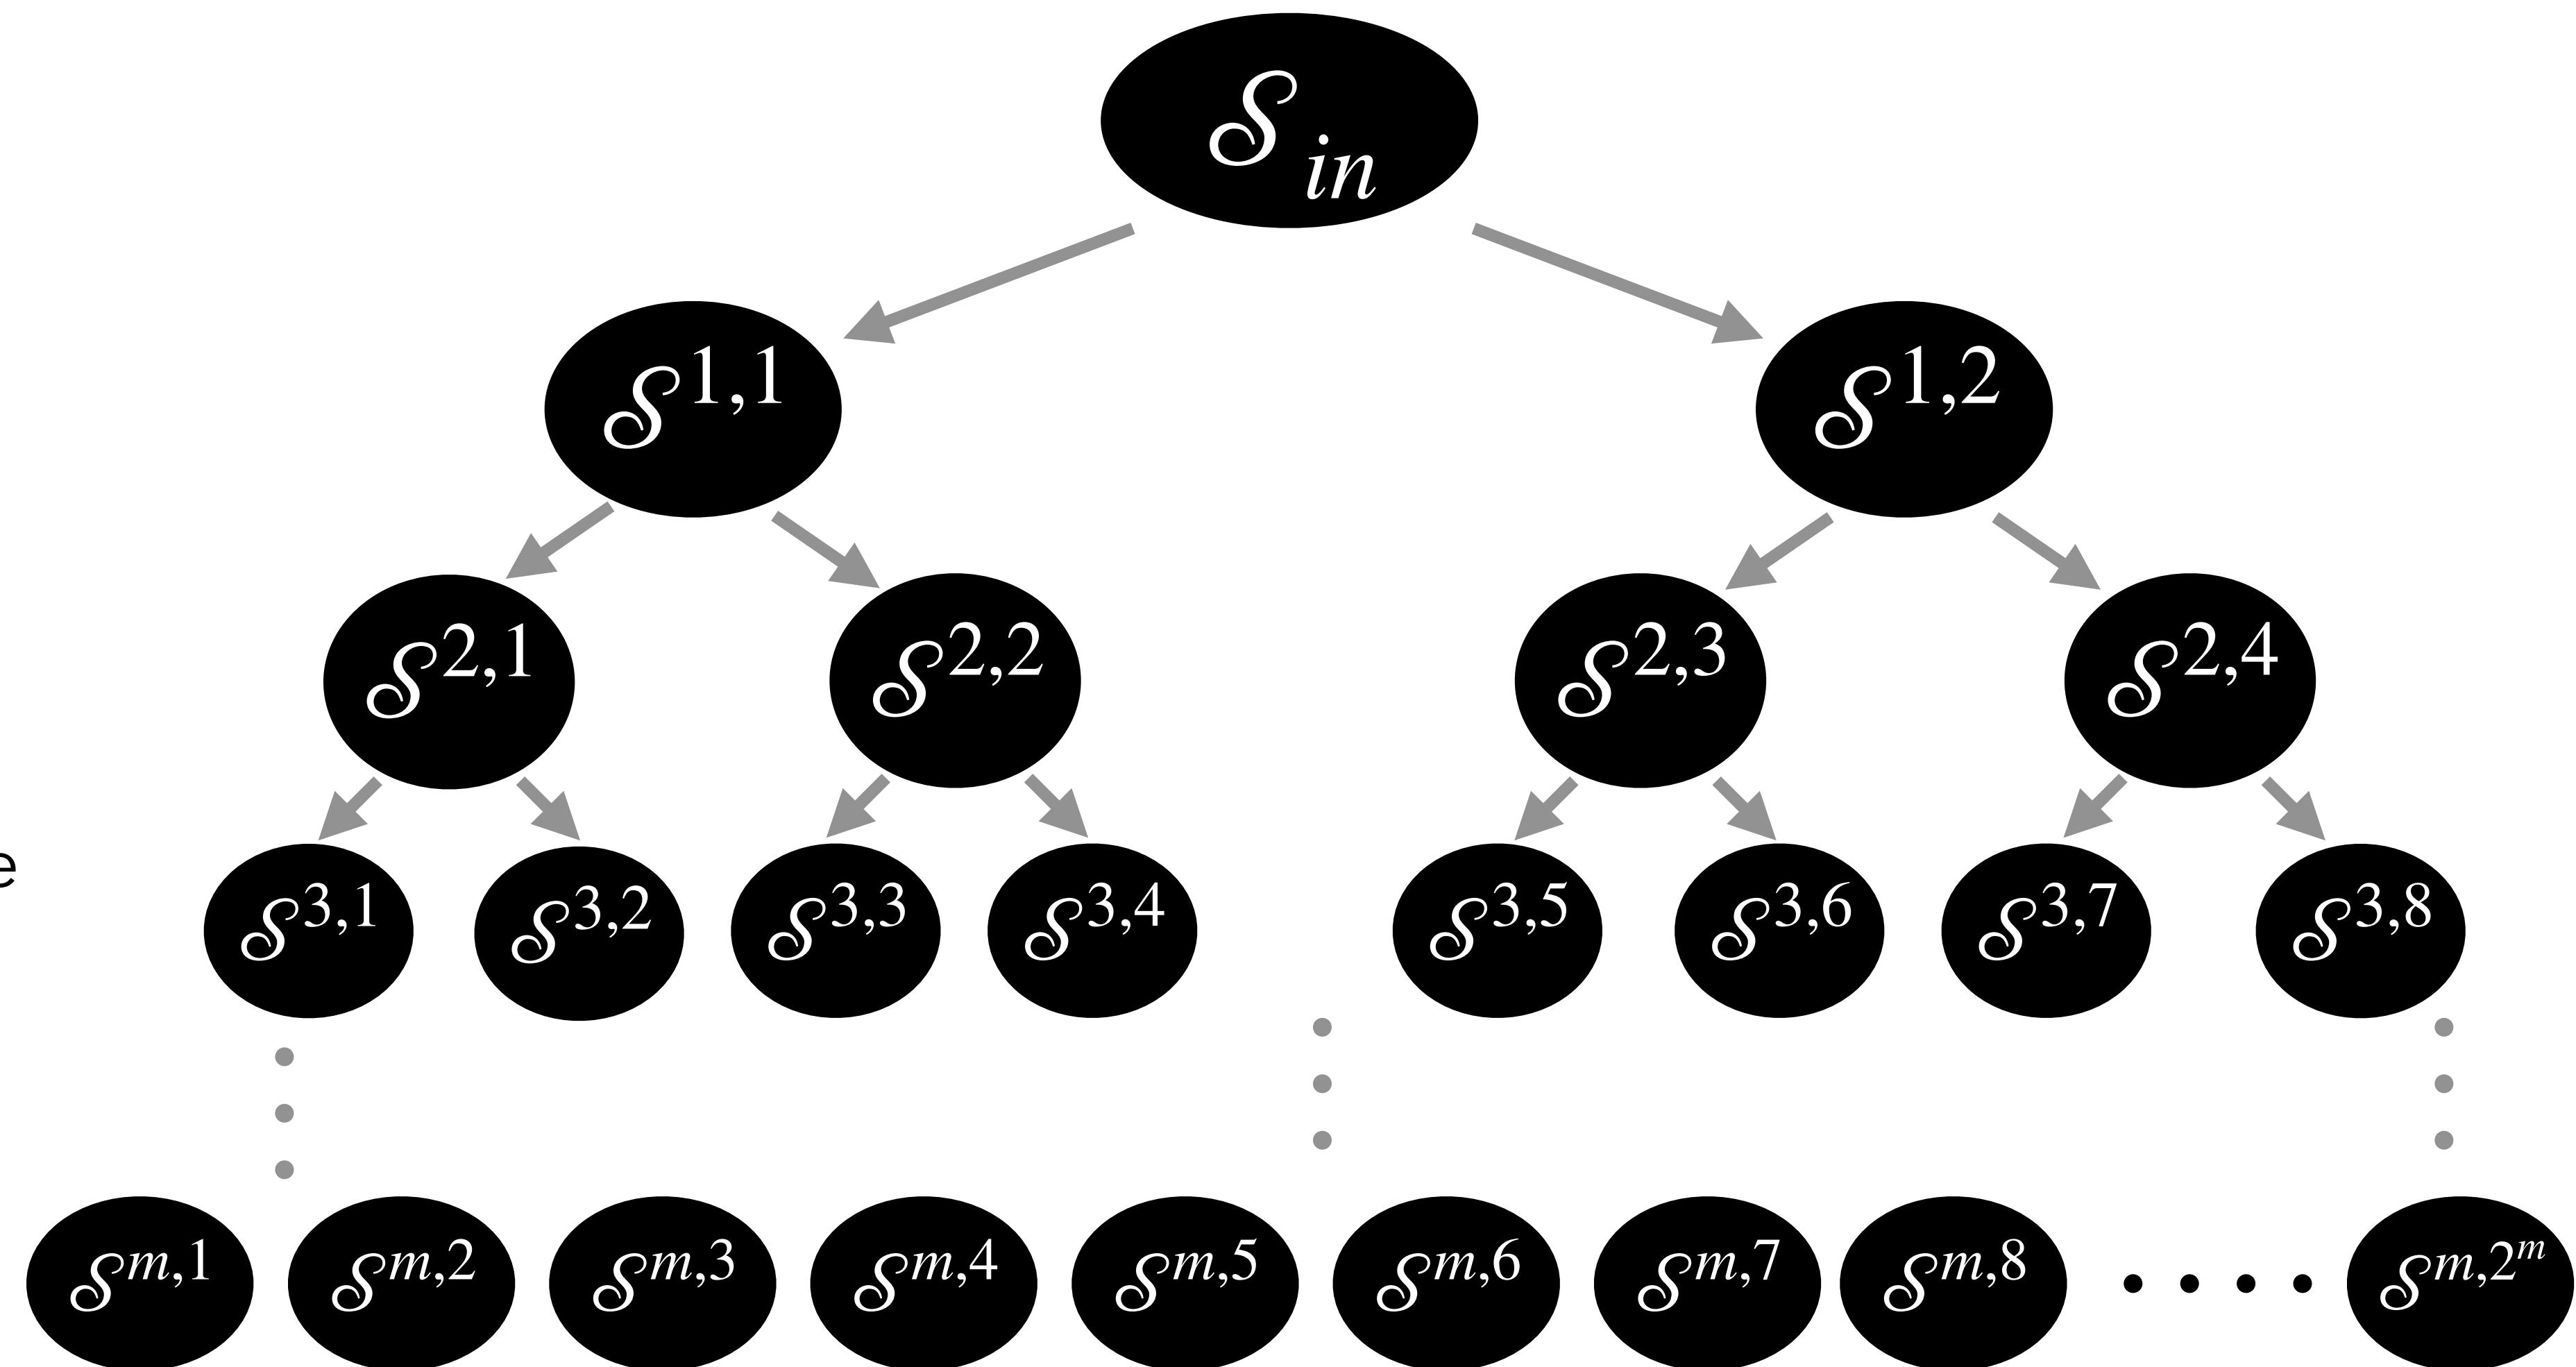
# KT: A two-staged algorithm

- **Input:** Kernel  $\mathbf{k}$ , input points  $\mathcal{S}_{in}$  of size  $n$ , thinning factor  $m$
- **KT-Split:**
  - Split  $\mathcal{S}_{in}$  into  $2^m$  balanced candidate coresets each of size  $\frac{n}{2^m}$
  - When  $m = \frac{1}{2} \log_2 n$ , we have  $\sqrt{n}$  coresets each of size  $\sqrt{n}$
- **KT-Swap:**
  - Pick the best candidate coreset that minimized  $\text{MMD}_{\mathbf{k}}$  to input
  - Iteratively refine each point in the selected coreset by swapping with the best alternative  $\mathcal{S}_{in}$  if it improves the MMD error

**Computation:**  $\mathcal{O}(n^2)$  kernel evaluations  
**Storage:**  $n \min(n, d)$

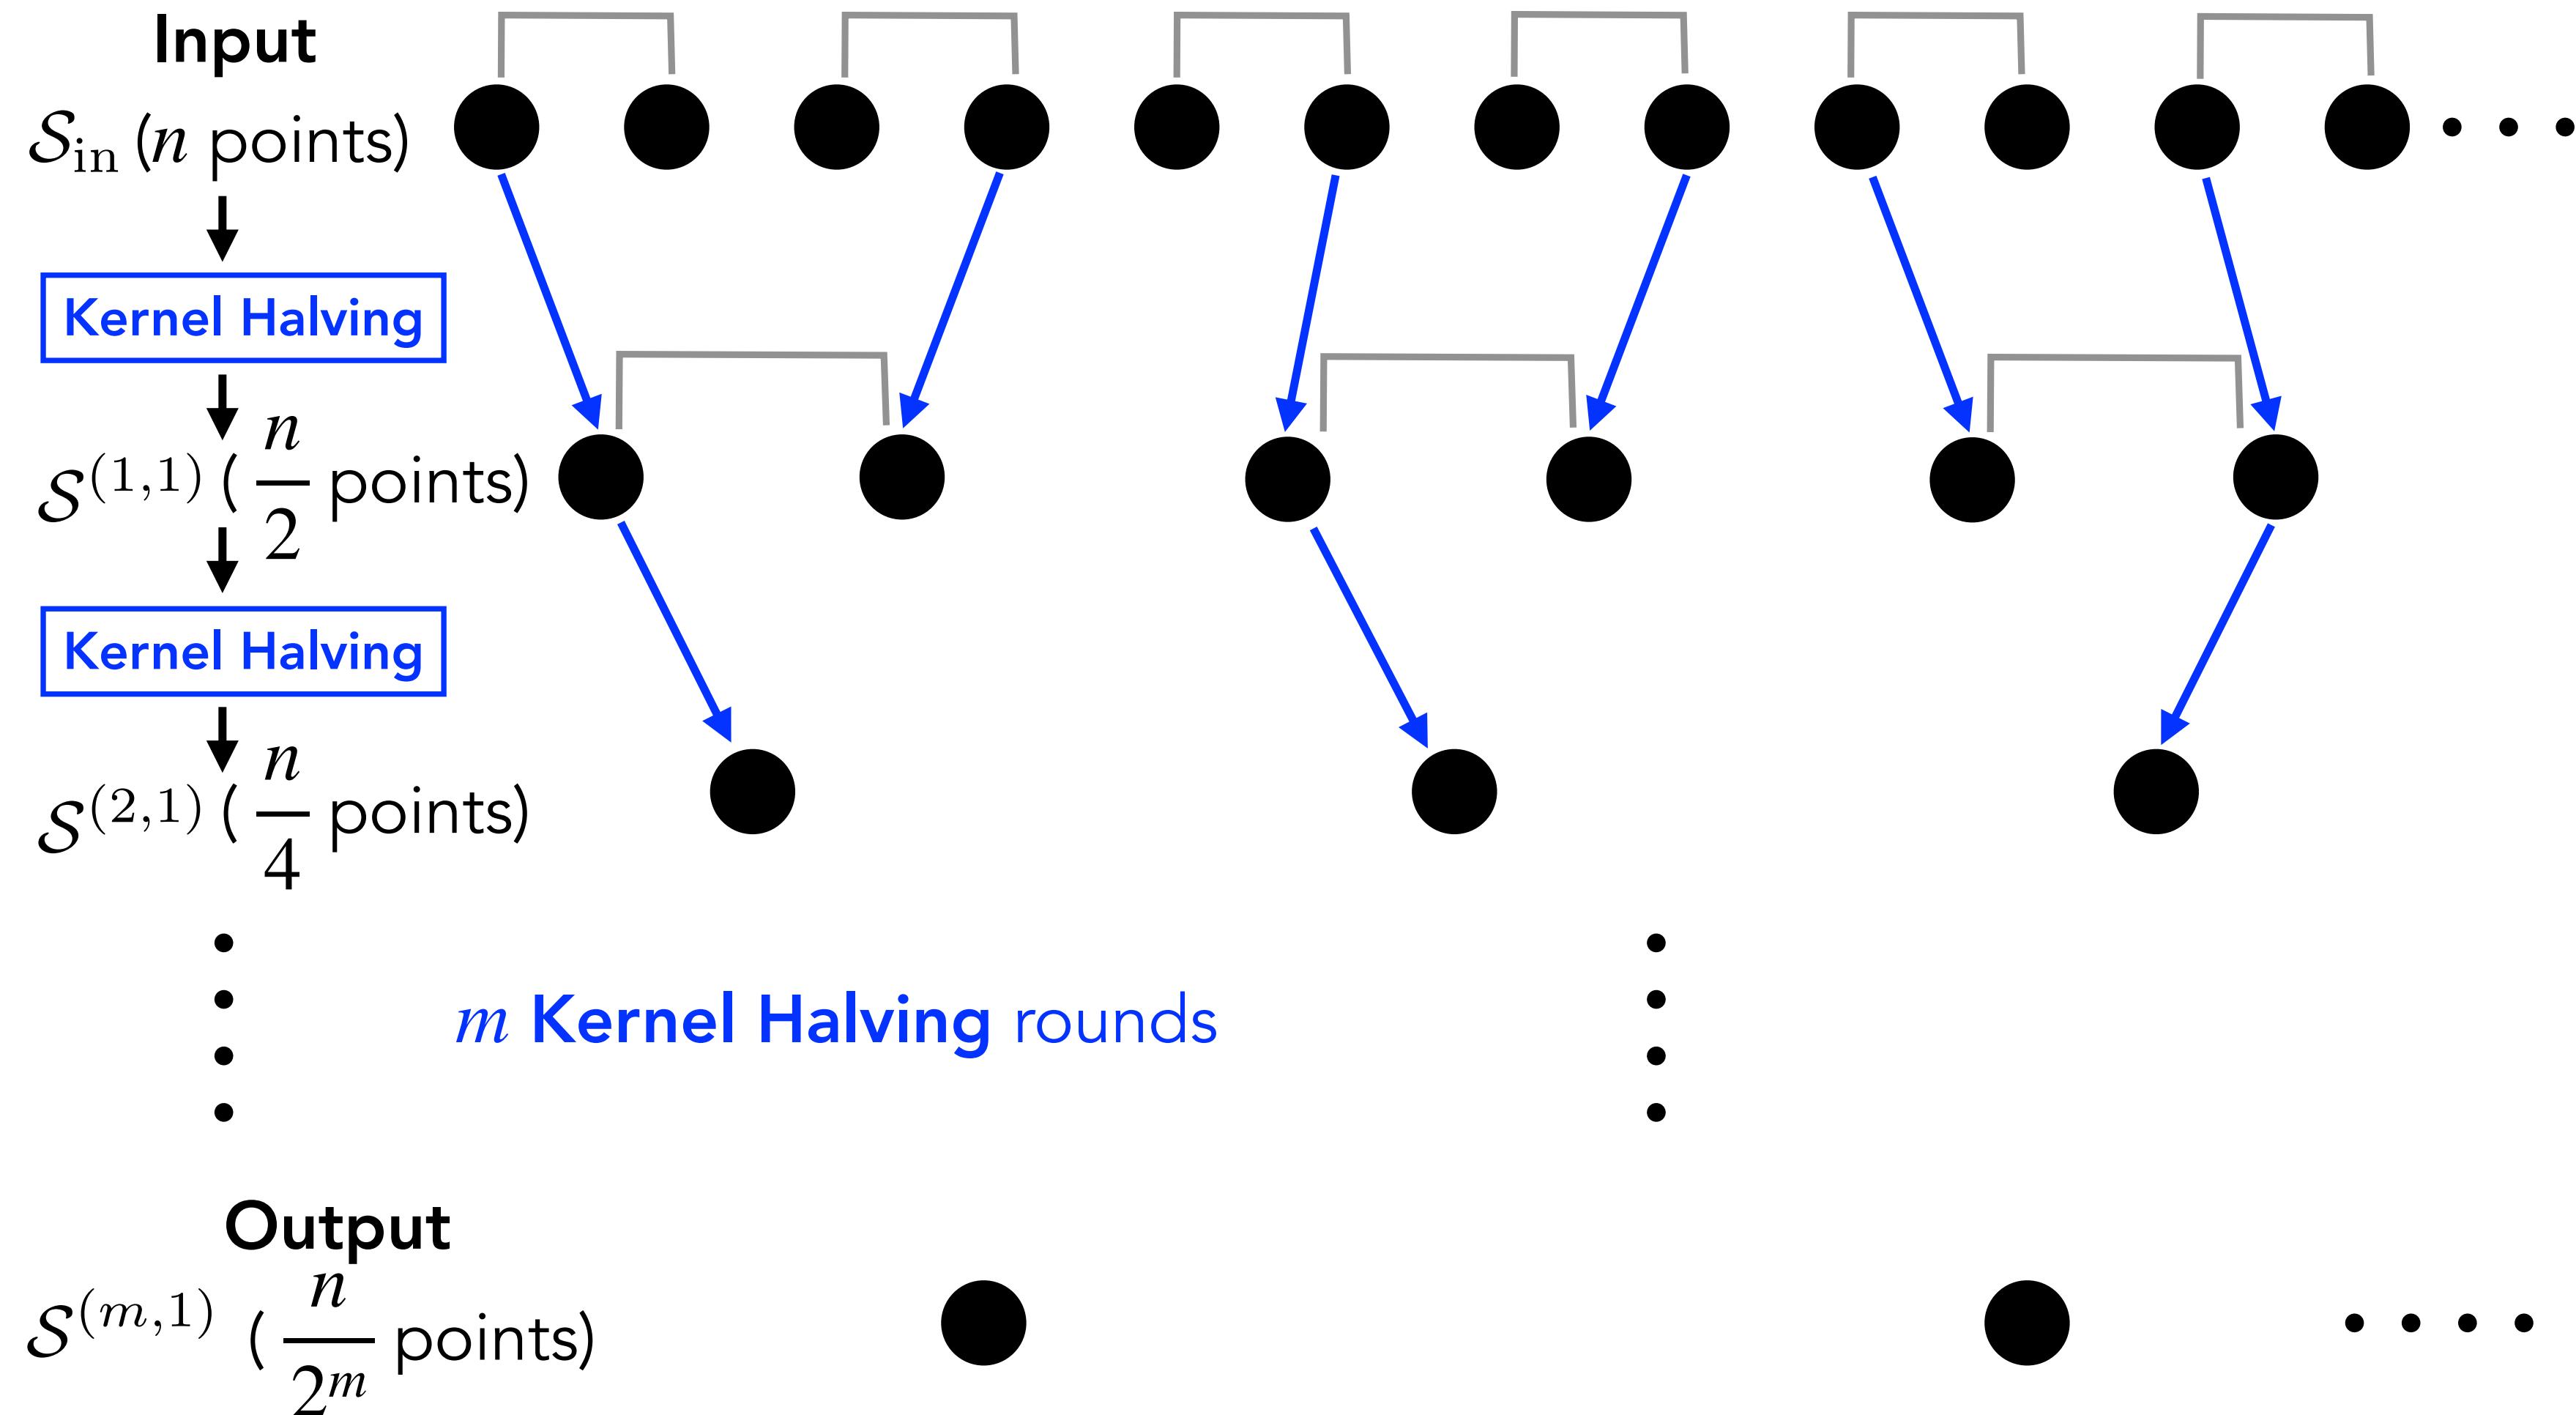
# KT-Split

- Repeated rounds of splitting the parent coresset in two equal-sized children coresets
  - Runs online, after seeing  $t$  input points, the bottom nodes have  $t/2^m$  points



# KT-Split

- One path on the tree is obtained by repeated **kernel halving**
- At each halving round, remaining points are paired, and one point is selected **non-uniformly** from each pair using a **new Hilbert space generalization of the self-balancing walk** of [Alweiss-Liu-Sawhney 2020]



---

**Algorithm 2:** Self-balancing Hilbert Walk

---

**Input:** sequence of functions  $(f_i)_{i=1}^n$  in Hilbert space  $\mathcal{H}$ , threshold sequence  $(\alpha_i)_{i=1}^n$

$\psi_0 \leftarrow 0 \in \mathcal{H}$

**for**  $i = 1, 2, \dots, n$  **do**

$\alpha_i \leftarrow \langle \psi_{i-1}, f_i \rangle_{\mathcal{H}}$  // Compute Hilbert space inner product

**if**  $|\alpha_i| > \alpha_i$ :

$\psi_i \leftarrow \psi_{i-1} - f_i \cdot \alpha_i / \alpha_i$

**else:**

$\eta_i \leftarrow 1$  with probability  $\frac{1}{2}(1 - \alpha_i / \alpha_i)$  and  $\eta_i \leftarrow -1$  otherwise

$\psi_i \leftarrow \psi_{i-1} + \eta_i f_i$

**end**

**return**  $\psi_n$ , combination of signed input functions

---

---

## Algorithm 2: Self-balancing Hilbert Walk

---

**Input:** sequence of functions  $(f_i)_{i=1}^n$  in Hilbert space  $\mathcal{H}$ , threshold sequence  $(\alpha_i)_{i=1}^n$

$\psi_0 \leftarrow 0 \in \mathcal{H}$

**for**  $i = 1, 2, \dots, n$  **do**

$\alpha_i \leftarrow \langle \psi_{i-1}, f_i \rangle_{\mathcal{H}}$  // Compute Hilbert space inner product

**if**  $|\alpha_i| > \alpha_i$ :

$\psi_i \leftarrow \psi_{i-1} - f_i \cdot \alpha_i / \alpha_i$

We choose  $\alpha_i$  such that this step **does not occur** with high probability

**else:**

$\eta_i \leftarrow 1$  with probability  $\frac{1}{2}(1 - \alpha_i / \alpha_i)$  and  $\eta_i \leftarrow -1$  otherwise

$\psi_i \leftarrow \psi_{i-1} + \eta_i f_i$

**end**

**return**  $\psi_n$ , combination of signed input functions

---

---

## Algorithm 2: Self-balancing Hilbert Walk

---

**Input:** sequence of functions  $(f_i)_{i=1}^n$  in Hilbert space  $\mathcal{H}$ , threshold sequence  $(\alpha_i)_{i=1}^n$

$\psi_0 \leftarrow 0 \in \mathcal{H}$

**for**  $i = 1, 2, \dots, n$  **do**

$\alpha_i \leftarrow \langle \psi_{i-1}, f_i \rangle_{\mathcal{H}}$  // Compute Hilbert space inner product

**if**  $|\alpha_i| > \alpha_i$ :

$\psi_i \leftarrow \psi_{i-1} - f_i \cdot \alpha_i / \alpha_i$  ← We choose  $\alpha_i$  such that this step **does not** occur with high probability

**else:**

$\eta_i \leftarrow 1$  with probability  $\frac{1}{2}(1 - \alpha_i / \alpha_i)$  and  $\eta_i \leftarrow -1$  otherwise

$\psi_i \leftarrow \psi_{i-1} + \eta_i f_i$

**end**

**return**  $\psi_n$ , combination of signed input functions

---

- Exact **Kernel halving**: When  $f_i = \mathbf{k}(2i, \cdot) - \mathbf{k}(x_{2i-1}, \cdot)$ , exactly half of input points ( $\mathcal{S}_{out}$ ) given  $-1$  sign after  $n/2$  steps

$$\frac{1}{n}\psi_n = \frac{1}{n} \sum_{x \in \mathcal{S}_{in}} \mathbf{k}(x, \cdot) - \frac{2}{n} \sum_{x \in \mathcal{S}_{out}} \mathbf{k}(x, \cdot) = \mathbb{P}_{in}\mathbf{k} - \mathbb{P}_{out}\mathbf{k}$$

---

**Algorithm 2:** Self-balancing Hilbert Walk

---

**Input:** sequence of functions  $(f_i)_{i=1}^n$  in Hilbert space  $\mathcal{H}$ , threshold sequence  $(\alpha_i)_{i=1}^n$

$\psi_0 \leftarrow 0 \in \mathcal{H}$

**for**  $i = 1, 2, \dots, n$  **do**

$\alpha_i \leftarrow \langle \psi_{i-1}, f_i \rangle_{\mathcal{H}}$  // Compute Hilbert space inner product

**if**  $|\alpha_i| > \alpha_i$ :

$\psi_i \leftarrow \psi_{i-1} - f_i \cdot \alpha_i / \alpha_i$

**else:**

$\eta_i \leftarrow 1$  with probability  $\frac{1}{2}(1 - \alpha_i / \alpha_i)$  and  $\eta_i \leftarrow -1$  otherwise

$\psi_i \leftarrow \psi_{i-1} + \eta_i f_i$

**end**

**return**  $\psi_n$ , combination of signed input functions

---

- **Balance:** If  $\mathbf{k}$  is a reproducing kernel, for all  $g \in \mathbb{H}_{\mathbf{k}'}$

$$\langle \psi_n, g \rangle_{\mathbf{k}} = \mathbb{P}_{in}g - \mathbb{P}_{out}g \text{ is } \mathcal{O}(n^{-1} \cdot \sqrt{\log n} \cdot \|g\|_{\mathbf{k}})\text{-sub-Gaussian}$$

If  $\eta_i$  were chosen i.i.d., the sub-Gaussian parameter is  $\Omega(n^{-1/2})$

# **Compress++**

# Compress++: A simple two-stage algorithm



---

**Algorithm 1:** COMPRESS

**Input:** halving algorithm HALVE, oversampling parameter  $g$ , point sequence  $S_{\text{in}}$  of size  $n$

**if**  $n = 4^g$  **then return**  $S_{in}$

**else**

Partition  $S_{\text{in}}$  into four arbitrary subsequences  $\{S_i\}_{i=1}^4$  each of size  $n/4$

**for**  $i = 1, 2, 3, 4$  **do**

```

 $\tilde{S}_i \leftarrow \text{COMPRESS}(S_i, \text{HALVE}, g)$  // return cores

```

end

$\tilde{\mathcal{S}} \leftarrow \text{CONCATENATE}(\tilde{\mathcal{S}}_1, \tilde{\mathcal{S}}_2, \tilde{\mathcal{S}}_3, \tilde{\mathcal{S}}_4)$  // coresets of size  $2 \cdot 2^g \cdot \sqrt{n}$

**return** HALVE( $\tilde{\mathcal{S}}$ ) // coresets of size  $2^g \sqrt{n}$

end

# $2^g$ thinning algorithm THIN

$\sqrt{n}$  points

# Compress++: A recursive strategy to reduce runtime for generic thinning algorithms

Thinning algorithm with  
 $n^2$  runtime and  
error  $e(n)$



$\sqrt{n}$ -thinning algorithm with  
 $n \log^3 n$  runtime and  
error  $4e(n)$

---

## Algorithm 1: COMPRESS

---

**Input:** halving algorithm HALVE, oversampling parameter  $g$ , point sequence  $S_{\text{in}}$  of size  $n$

**if**  $n = 4^g$  **then return**  $S_{\text{in}}$

**else**

    Partition  $S_{\text{in}}$  into four arbitrary subsequences  $\{S_i\}_{i=1}^4$  each of size  $n/4$

**for**  $i = 1, 2, 3, 4$  **do**

$\tilde{S}_i \leftarrow \text{COMPRESS}(S_i, \text{HALVE}, g)$  // return coresets of size  $2^g \cdot \sqrt{\frac{n}{4}}$

**end**

$\tilde{S} \leftarrow \text{CONCATENATE}(\tilde{S}_1, \tilde{S}_2, \tilde{S}_3, \tilde{S}_4)$  // coreset of size  $2 \cdot 2^g \cdot \sqrt{n}$

**return**  $\text{HALVE}(\tilde{S})$  // coreset of size  $2^g \sqrt{n}$

**end**

---

# Compress++: Informal guarantee

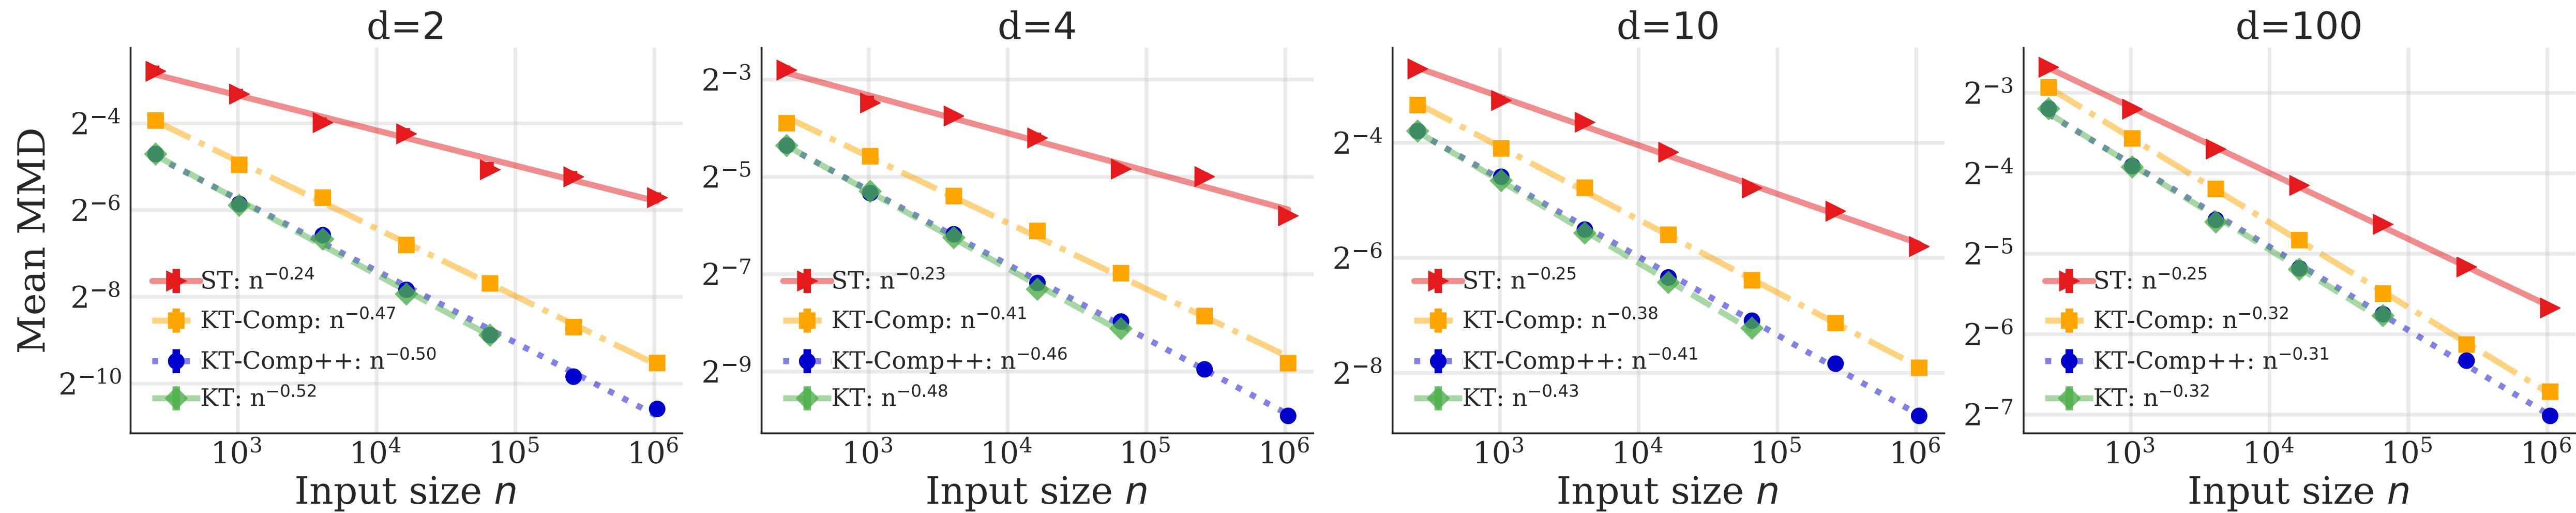
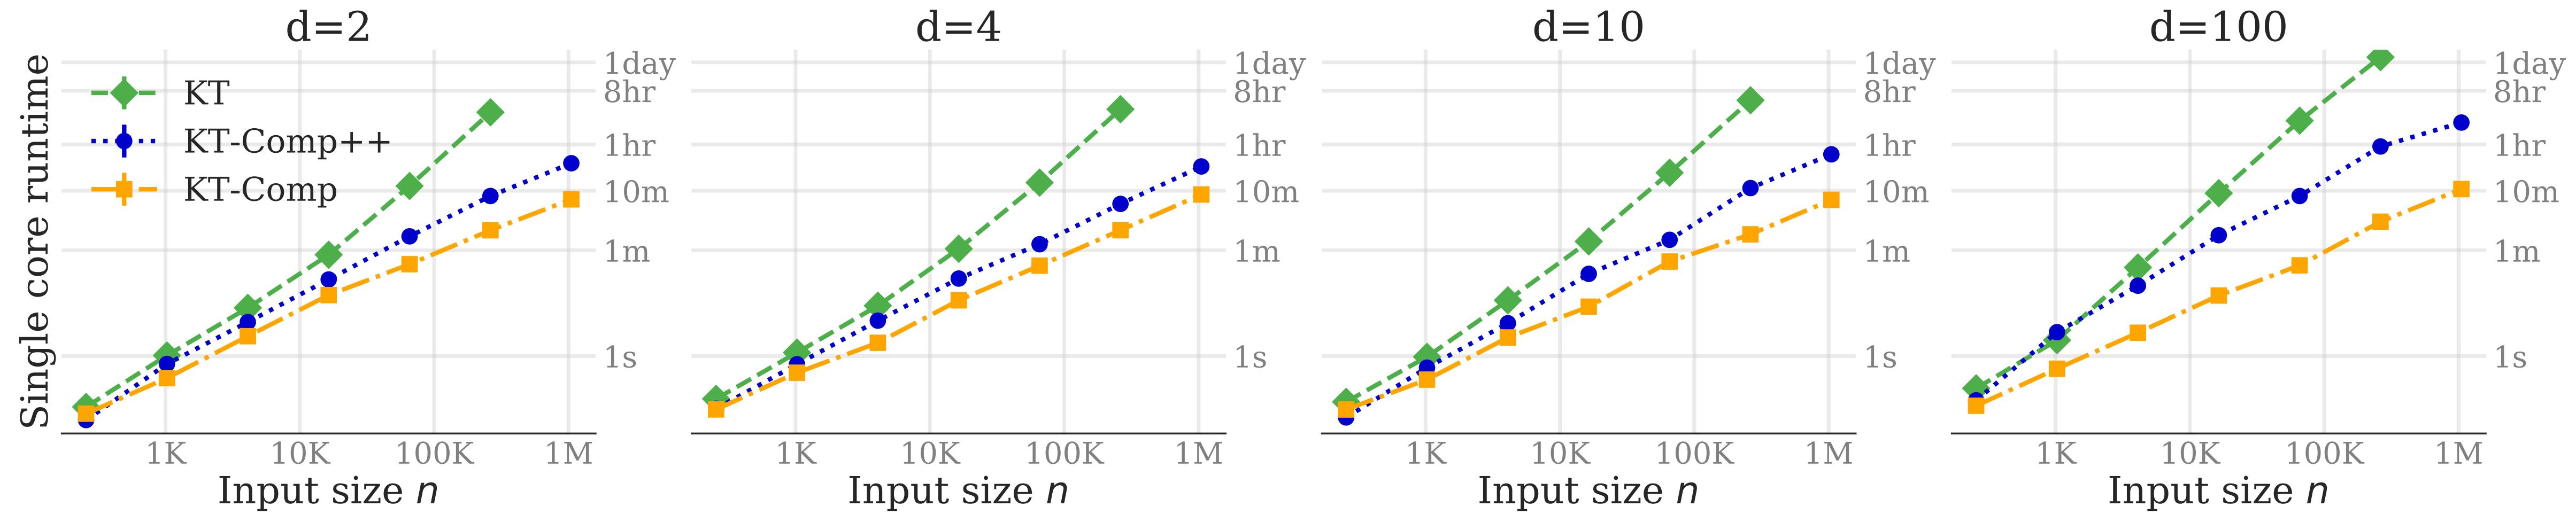
- Under some mild conditions, with  $g = \log \log n + 1$ , we have
- **Error inflation by at most 4:**

If  $\text{MMD}_{\mathbf{k}}(\mathbb{P}_n, \mathbb{P}_{HALVE}) \sim e_1(n)$  and  $\text{MMD}_{\mathbf{k}}(\mathbb{P}_n, \mathbb{P}_{THIN}) \sim e_2(n)$  then  
 $\text{MMD}_{\mathbf{k}}(\mathbb{P}_n, \mathbb{P}_{Compress++}) \sim 4 \max(e_1(n), e_2(n))$
- **Quadratic reduction in runtime:**

If runtime of HALVE and THIN with  $n$  points is  $\mathcal{O}(n^\tau)$  then the runtime of Compress++ with  $n$  points is  $\mathcal{O}(n^{\tau/2})$  if  $\tau > 2$  and  $\mathcal{O}(n \log^3 n)$  if  $\tau = 2$ .

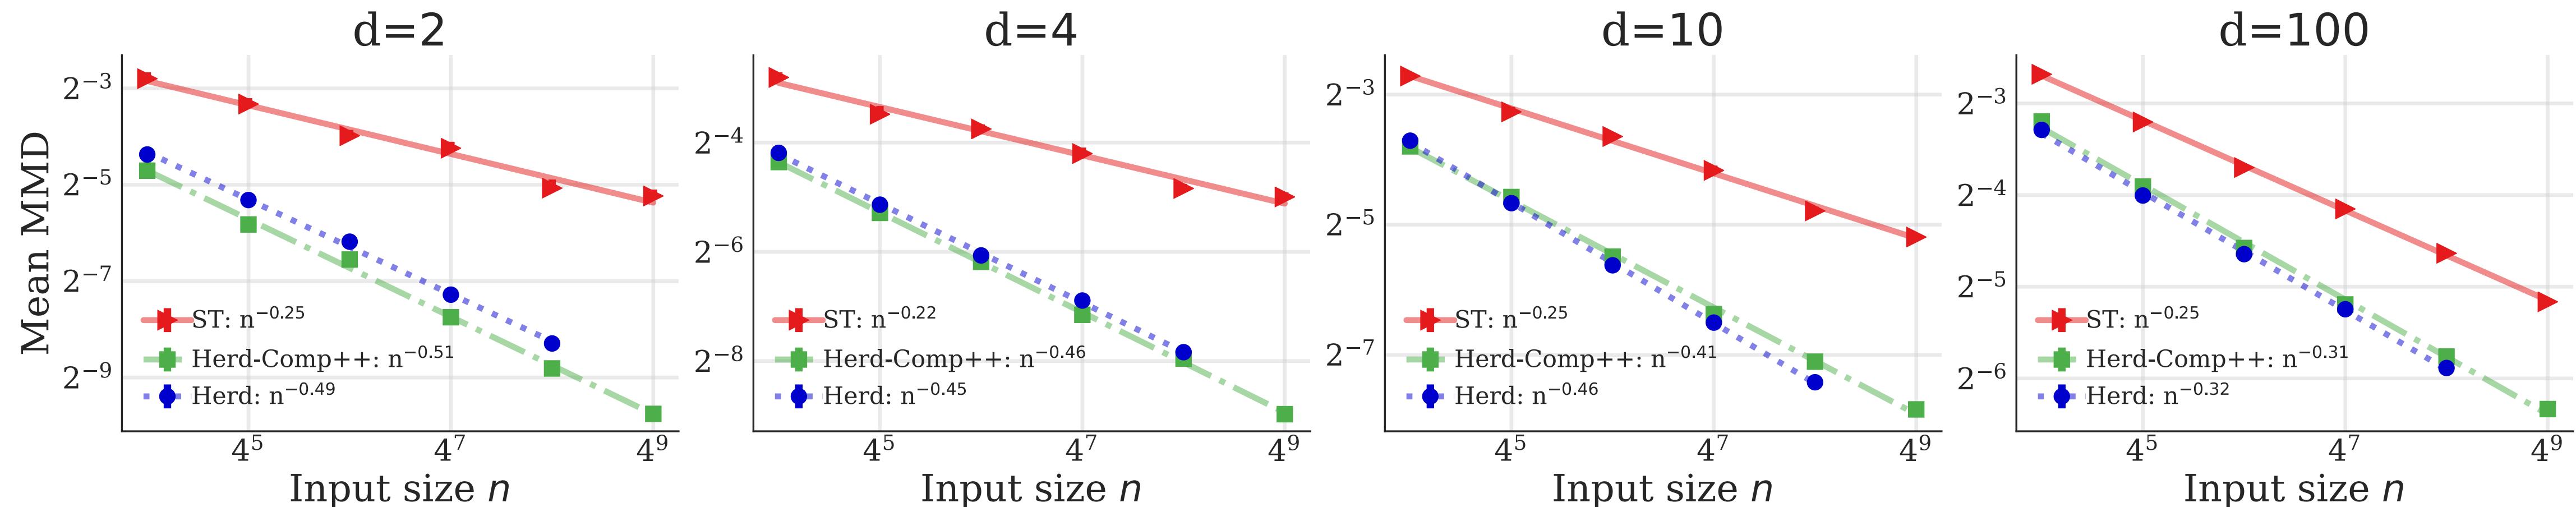
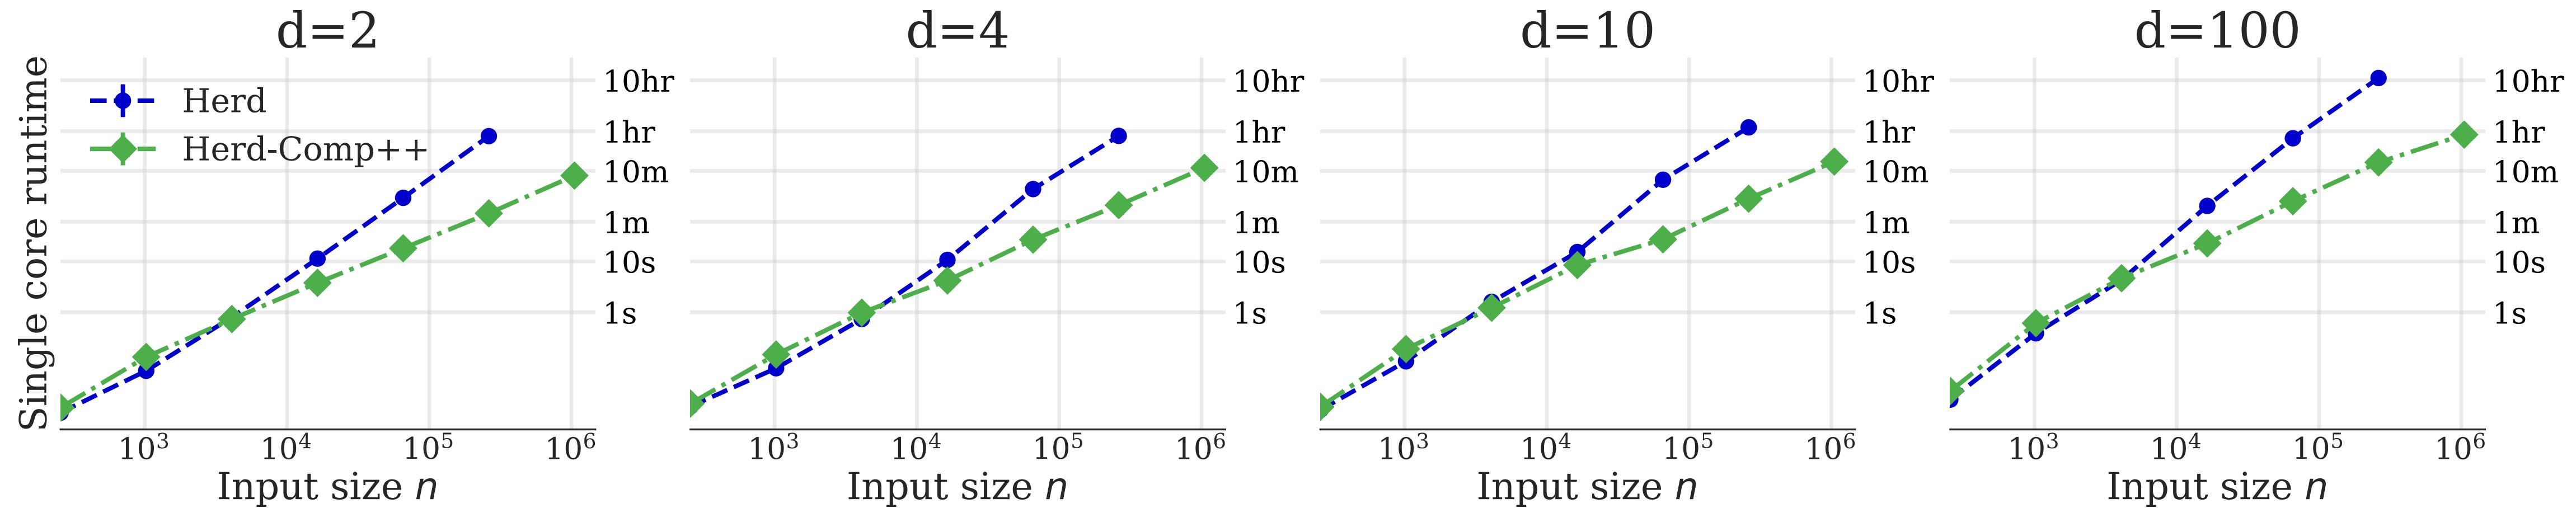
# KT vs Compress ( $g = 0$ ) vs Compress++ ( $g = 4$ )

The input algorithms Halve and Thin to Compress++ are derived from KT

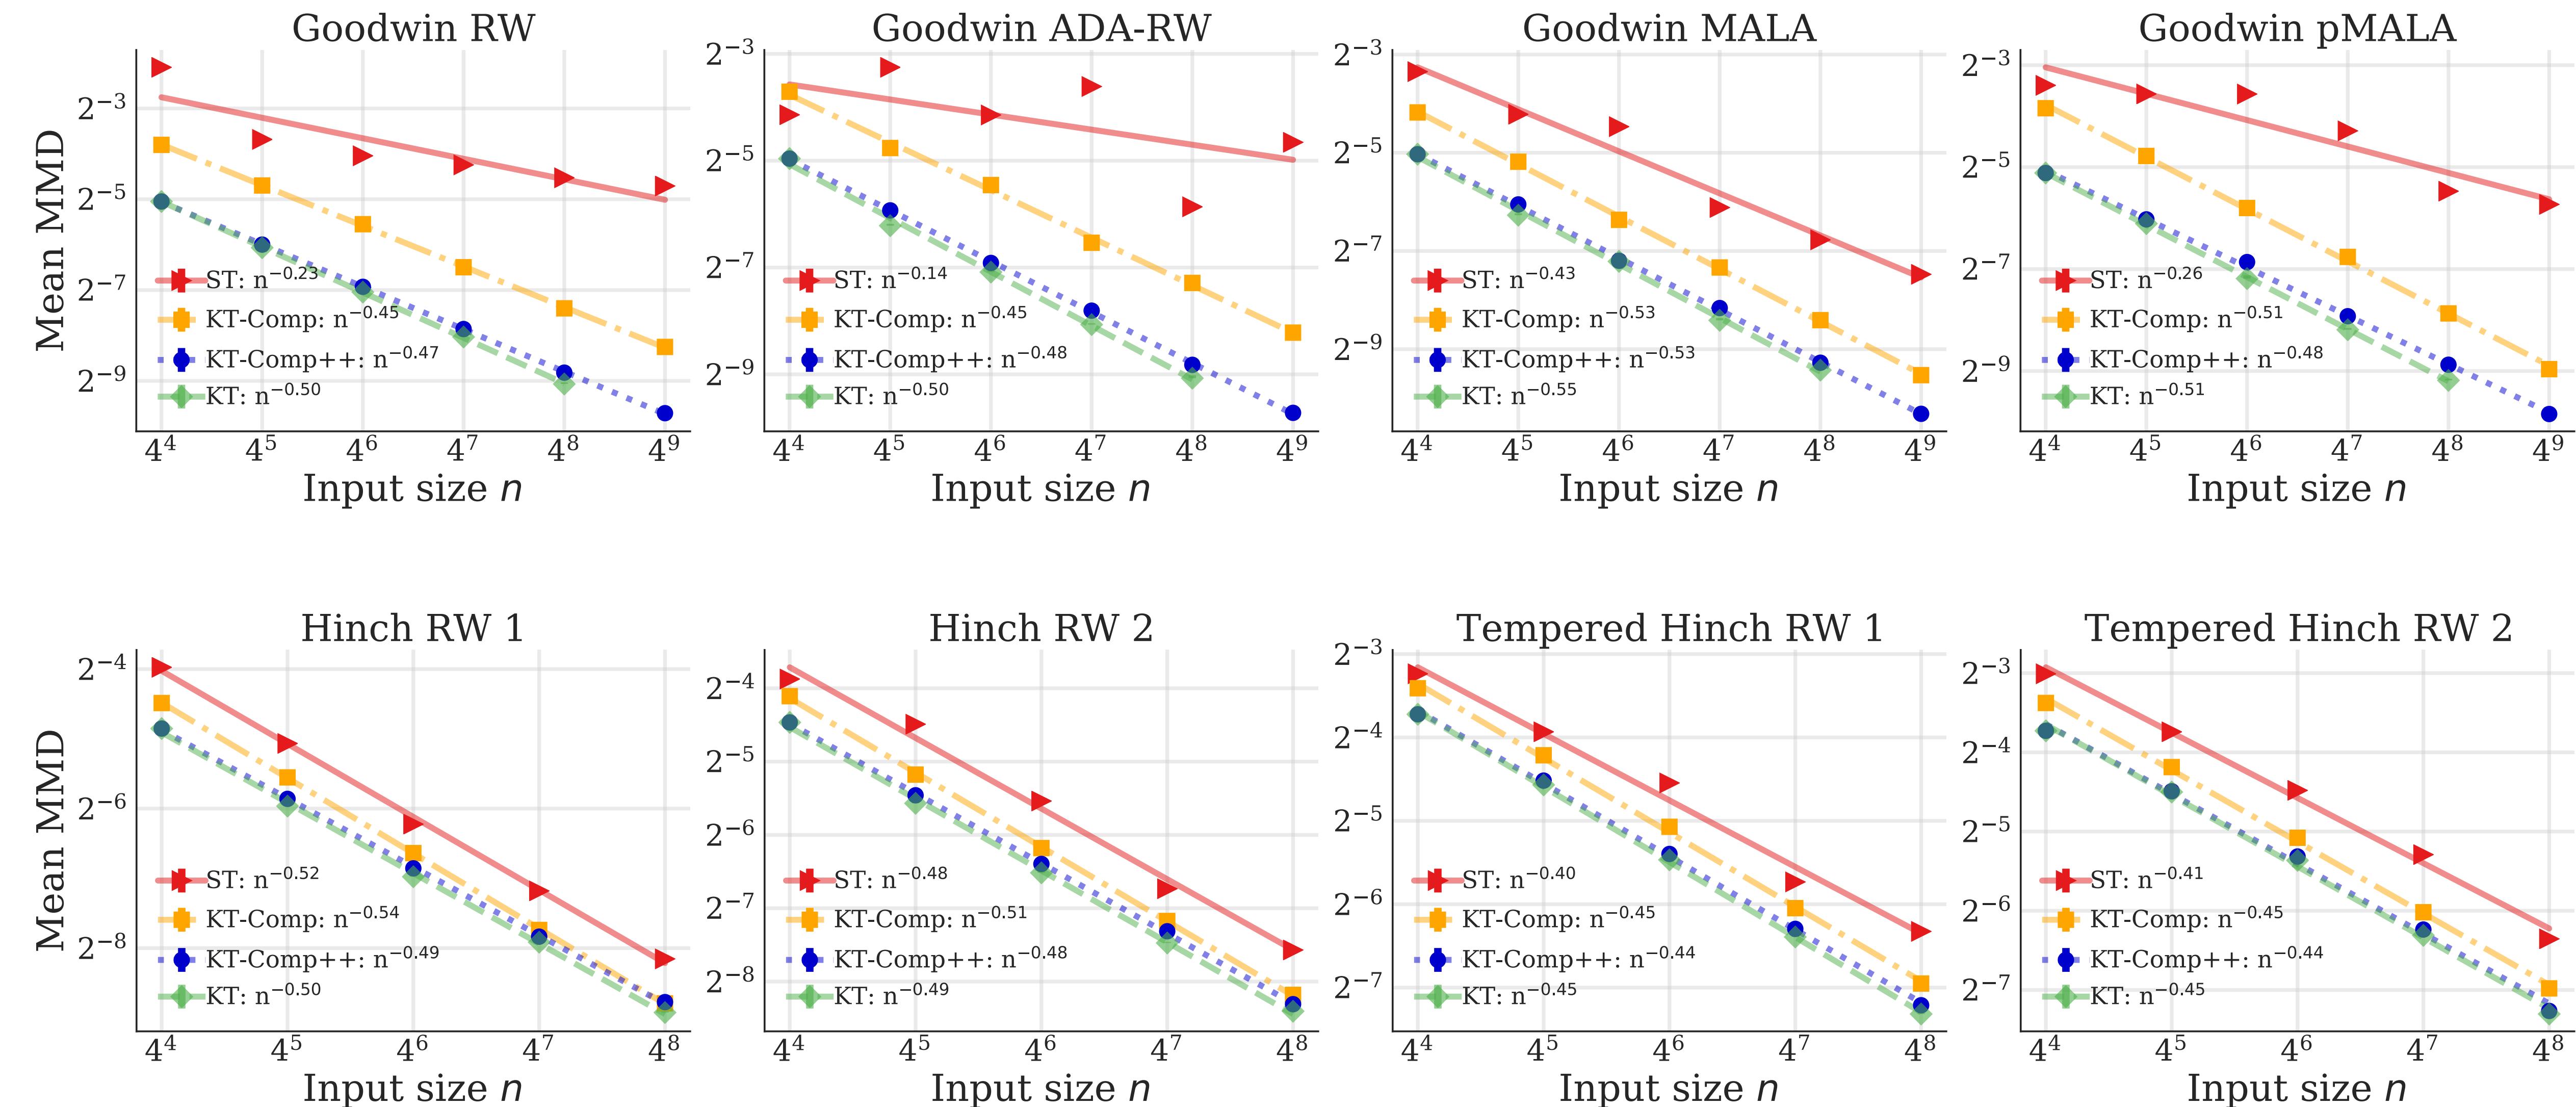


# Compress++: A meta strategy to reduce runtime

Results for Compress++ with kernel herding (Herd)



# MCMC Experiments



# Lower bounds

# Lower bounds

- For smooth kernels, there exists a target  $\mathbb{P}$ , such that a coresnet of size  $\sqrt{n}$  suffers an MMD error of  $\min(\sqrt{\frac{d}{n}}, n^{-1/4})$ . [Philips and Tai 2020]
- For characteristic kernels, there exists a target  $\mathbb{P}$ , such that any estimator based on  $n$  i.i.d. input points must suffer at least  $n^{-1/2}$  MMD error. [Tolstikhin et al. 2017]

**Both bounds apply to Gaussian and Matérn kernels**

# MCMC Details

# MCMC experiments: Differential equation models

**Dimension  $d = 4$**

1. Lotka-Volterra model

oscillatory enzymatic control, [1925, 1926]

2. Goodwin model

oscillatory predator-prey evolution, [1965]

# MCMC experiments: Differential equation models

**Dimension  $d = 4$**

1. Lotka-Volterra model

oscillatory enzymatic control, [1925, 1926]

2. Goodwin model

oscillatory predator-prey evolution, [1965]

X

1. Posterior

# MCMC experiments: Differential equation models

**Dimension  $d = 4$**

1. Lotka-Volterra model

oscillatory enzymatic control, [1925, 1926]

2. Goodwin model

oscillatory predator-prey evolution, [1965]

X

1. Posterior

X

1. Random walk (RW)  
[Metropolis et al. 1953, Hastings 1970]
2. Adaptive random walk (adaRW)  
[Haario et al. 1999]
3. Metropolis adjusted Langevin algorithm (MALA) [Roberts et al. 1996]
4. Preconditioned-MALA (pMALA)  
[Girolami et al. 2011]

# MCMC experiments: Differential equation models

**Dimension  $d = 4$**

1. Lotka-Volterra model

oscillatory enzymatic control, [1925, 1926]

2. Goodwin model

oscillatory predator-prey evolution, [1965]

X

1. Posterior

X

1. Random walk (RW)  
[Metropolis et al. 1953, Hastings 1970]
2. Adaptive random walk (adaRW)  
[Haario et al. 1999]
3. Metropolis adjusted Langevin algorithm (MALA) [Roberts et al. 1996]
4. Preconditioned-MALA (pMALA)  
[Girolami et al. 2011]

**Dimension  $d = 38$**

3. Hinch calcium signal model

[Hinch-Greenstein-Tanskanen-Xu-Winslow, 2004]

# MCMC experiments: Differential equation models

**Dimension  $d = 4$**

1. Lotka-Volterra model

oscillatory enzymatic control, [1925, 1926]

2. Goodwin model

oscillatory predator-prey evolution, [1965]

X

1. Posterior

X

1. Random walk (RW)  
[Metropolis et al. 1953, Hastings 1970]
2. Adaptive random walk (adaRW)  
[Haario et al. 1999]
3. Metropolis adjusted Langevin algorithm (MALA) [Roberts et al. 1996]
4. Preconditioned-MALA (pMALA)  
[Girolami et al. 2011]

**Dimension  $d = 38$**

3. Hinch calcium signal model

[Hinch-Greenstein-Tanskanen-Xu-Winslow, 2004]

X

1. Posterior  
2. Tempered posterior

X

1. Random walk (RW) - run 1
2. Random walk (RW) - run 2

# MCMC experiments: Differential equation models

**Dimension  $d = 4$**

1. Lotka-Volterra model

oscillatory enzymatic control, [1925, 1926]

2. Goodwin model

oscillatory predator-prey evolution, [1965]

X

1. Posterior

X

1. Random walk (RW)  
[Metropolis et al. 1953, Hastings 1970]
2. Adaptive random walk (adaRW)  
[Haario et al. 1999]
3. Metropolis adjusted Langevin algorithm (MALA) [Roberts et al. 1996]
4. Preconditioned-MALA (pMALA)  
[Girolami et al. 2011]

**Dimension  $d = 38$**

3. Hinch calcium signal model

[Hinch-Greenstein-Tanskanen-Xu-Winslow, 2004]

X

1. Posterior  
2. Tempered posterior

X

1. Random walk (RW) - run 1
2. Random walk (RW) - run 2

**For KT, we use Gaussian kernel, and chose its bandwidth via median heuristic** [Garreau et al. 2017]

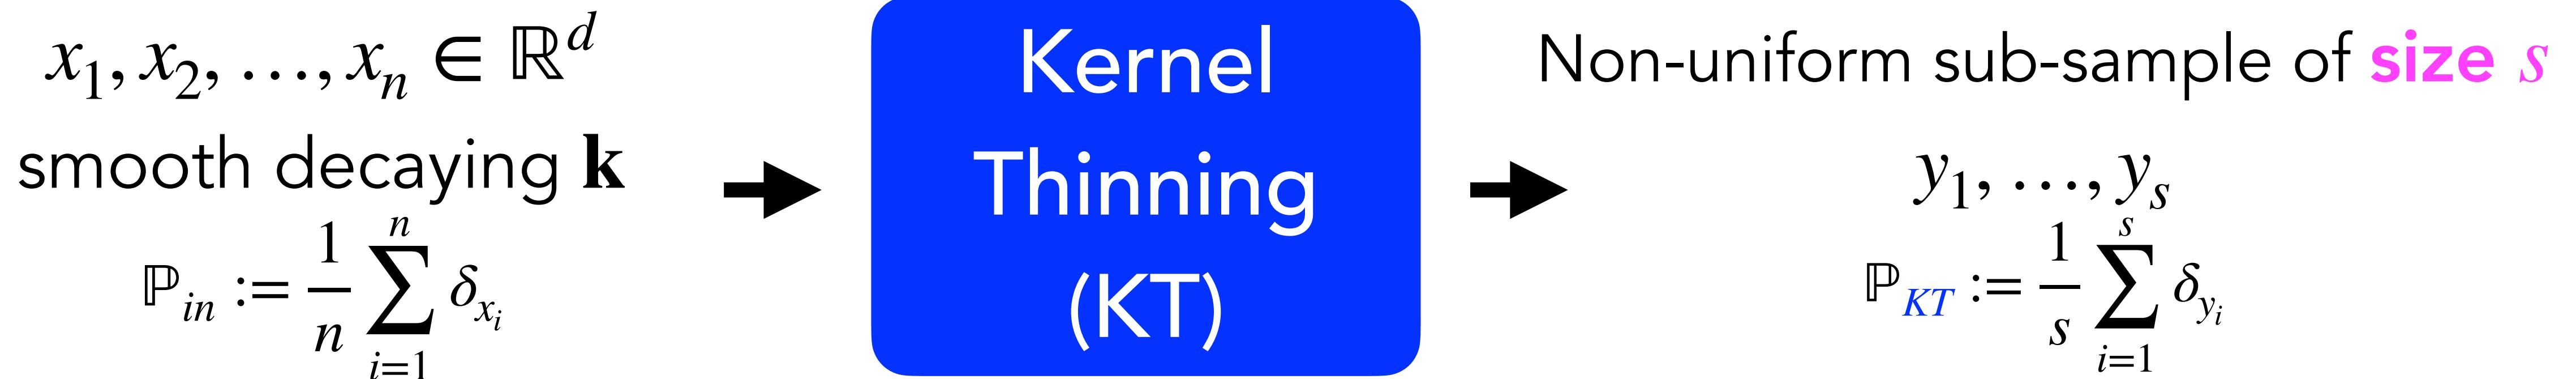
# Exact constants for KT result

# Generalized kernel thinning

	Root KT	Target KT	KT+ [Best of both worlds]
KT-Split kernel	$\mathbf{k}_{rt}$	$\mathbf{k}$	$\mathbf{k} + \mathbf{k}_{\alpha-rt}$
Single-function error	Same as MMD error	$\sqrt{\frac{\log n}{n}}$ For arbitrary $\mathbf{k}$ on arbitrary domain	$\sqrt{\frac{\log n}{n}}$
MMD error	See <u>slide</u>	$\sqrt{\frac{\log^{ad+b} n}{n}}$ Analytic $\mathbf{k}$ $\sqrt{\frac{n^{d/m}}{n}}$ $m$ -times differentiable $\mathbf{k}$	Min(Target KT Error, $\alpha$ -Root KT)

$$\mathbf{k}(x, y) = \int \mathbf{k}_{rt}(x, z) \mathbf{k}_{rt}(z, y) dz \quad \& \quad \mathbf{k}_{\alpha-rt} = \widehat{(\mathbf{k})}^\alpha \text{ where } \widehat{\cdot} \text{ denotes Fourier transform}$$

# Target KT or KT+: Better than Monte Carlo rate

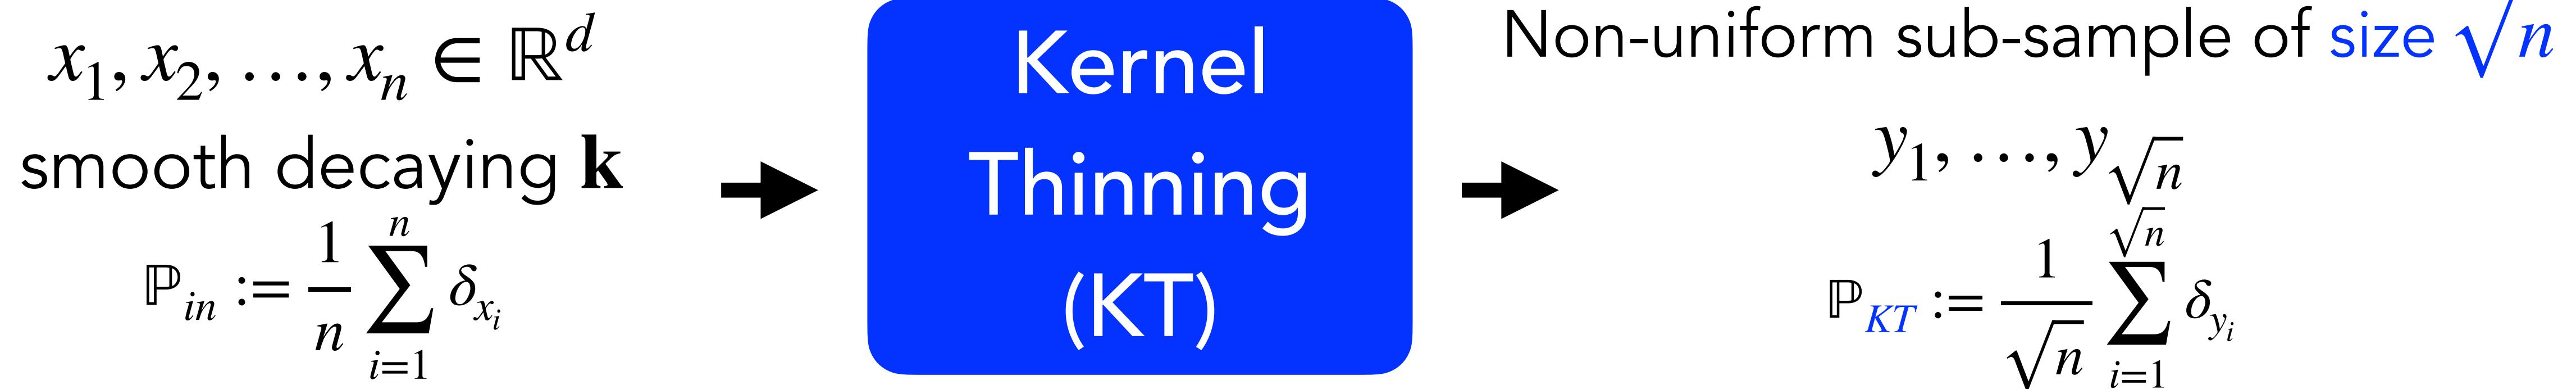


For any fixed  $g \in \mathbb{H}_k$ , with probability  $1 - \delta$  over the randomness in KT, we have

$$|\mathbb{P}_{in}g - \mathbb{P}_{KT}g| \leq \frac{1}{s} \cdot \|g\|_k \sqrt{\frac{8}{3} \|\mathbf{k}\|_\infty \log\left(\frac{4}{\delta}\right) \log\left(\frac{6s \log(n/s)}{\delta}\right)}$$

Much faster than the Monte Carlo rate for standard/uniform thinning  $\mathcal{O}\left(\frac{1}{\sqrt{s}}\right)$

# Target KT or KT+: Better than Monte Carlo rate



For any fixed  $g \in \mathbb{H}_k$ , with probability  $1 - \delta$  over the randomness in KT, we have

$$|\mathbb{P}_{in}g - \mathbb{P}_{KT}g| \leq \frac{1}{\sqrt{n}} \cdot \|g\|_k \sqrt{\frac{8}{3} \|\mathbf{k}\|_\infty \log\left(\frac{4}{\delta}\right) \log\left(\frac{6\sqrt{n} \log \sqrt{n}}{\delta}\right)}$$

Much faster than the Monte Carlo rate for standard/uniform thinning  $\mathcal{O}\left(\frac{1}{n^{1/4}}\right)$

# Properties of MMD

- Maximum mean discrepancy (MMD) = worst-case integration discrepancy between two distributions over a class of real-valued test functions

$$\text{MMD}_{\mathbf{k}}(\mathbb{P}_{in}, \mathbb{P}_{out}) = \sup_{\|g\|_{\mathbf{k}} \leq 1} |\mathbb{P}_{in}g - \mathbb{P}_{out}g|$$

[Gretton-Borgwardt-Rasch-Schölkopf-Smola, 2012]

- **MMD metrizes convergence in distribution** for popular infinite-dimensional kernels like Gaussian, Matern, IMQ, B-spline

[Simon-Gabriel-Barp-Schölkopf-Mackey, 2020]

Target KT MMD rates:  $\sqrt{n}$  points with  $\widetilde{O}(n^{-1/2})$  error

- More generally

$$\widetilde{O}\left(\frac{1}{\sqrt{n}}\right) \text{ error rate for analytic kernels}$$

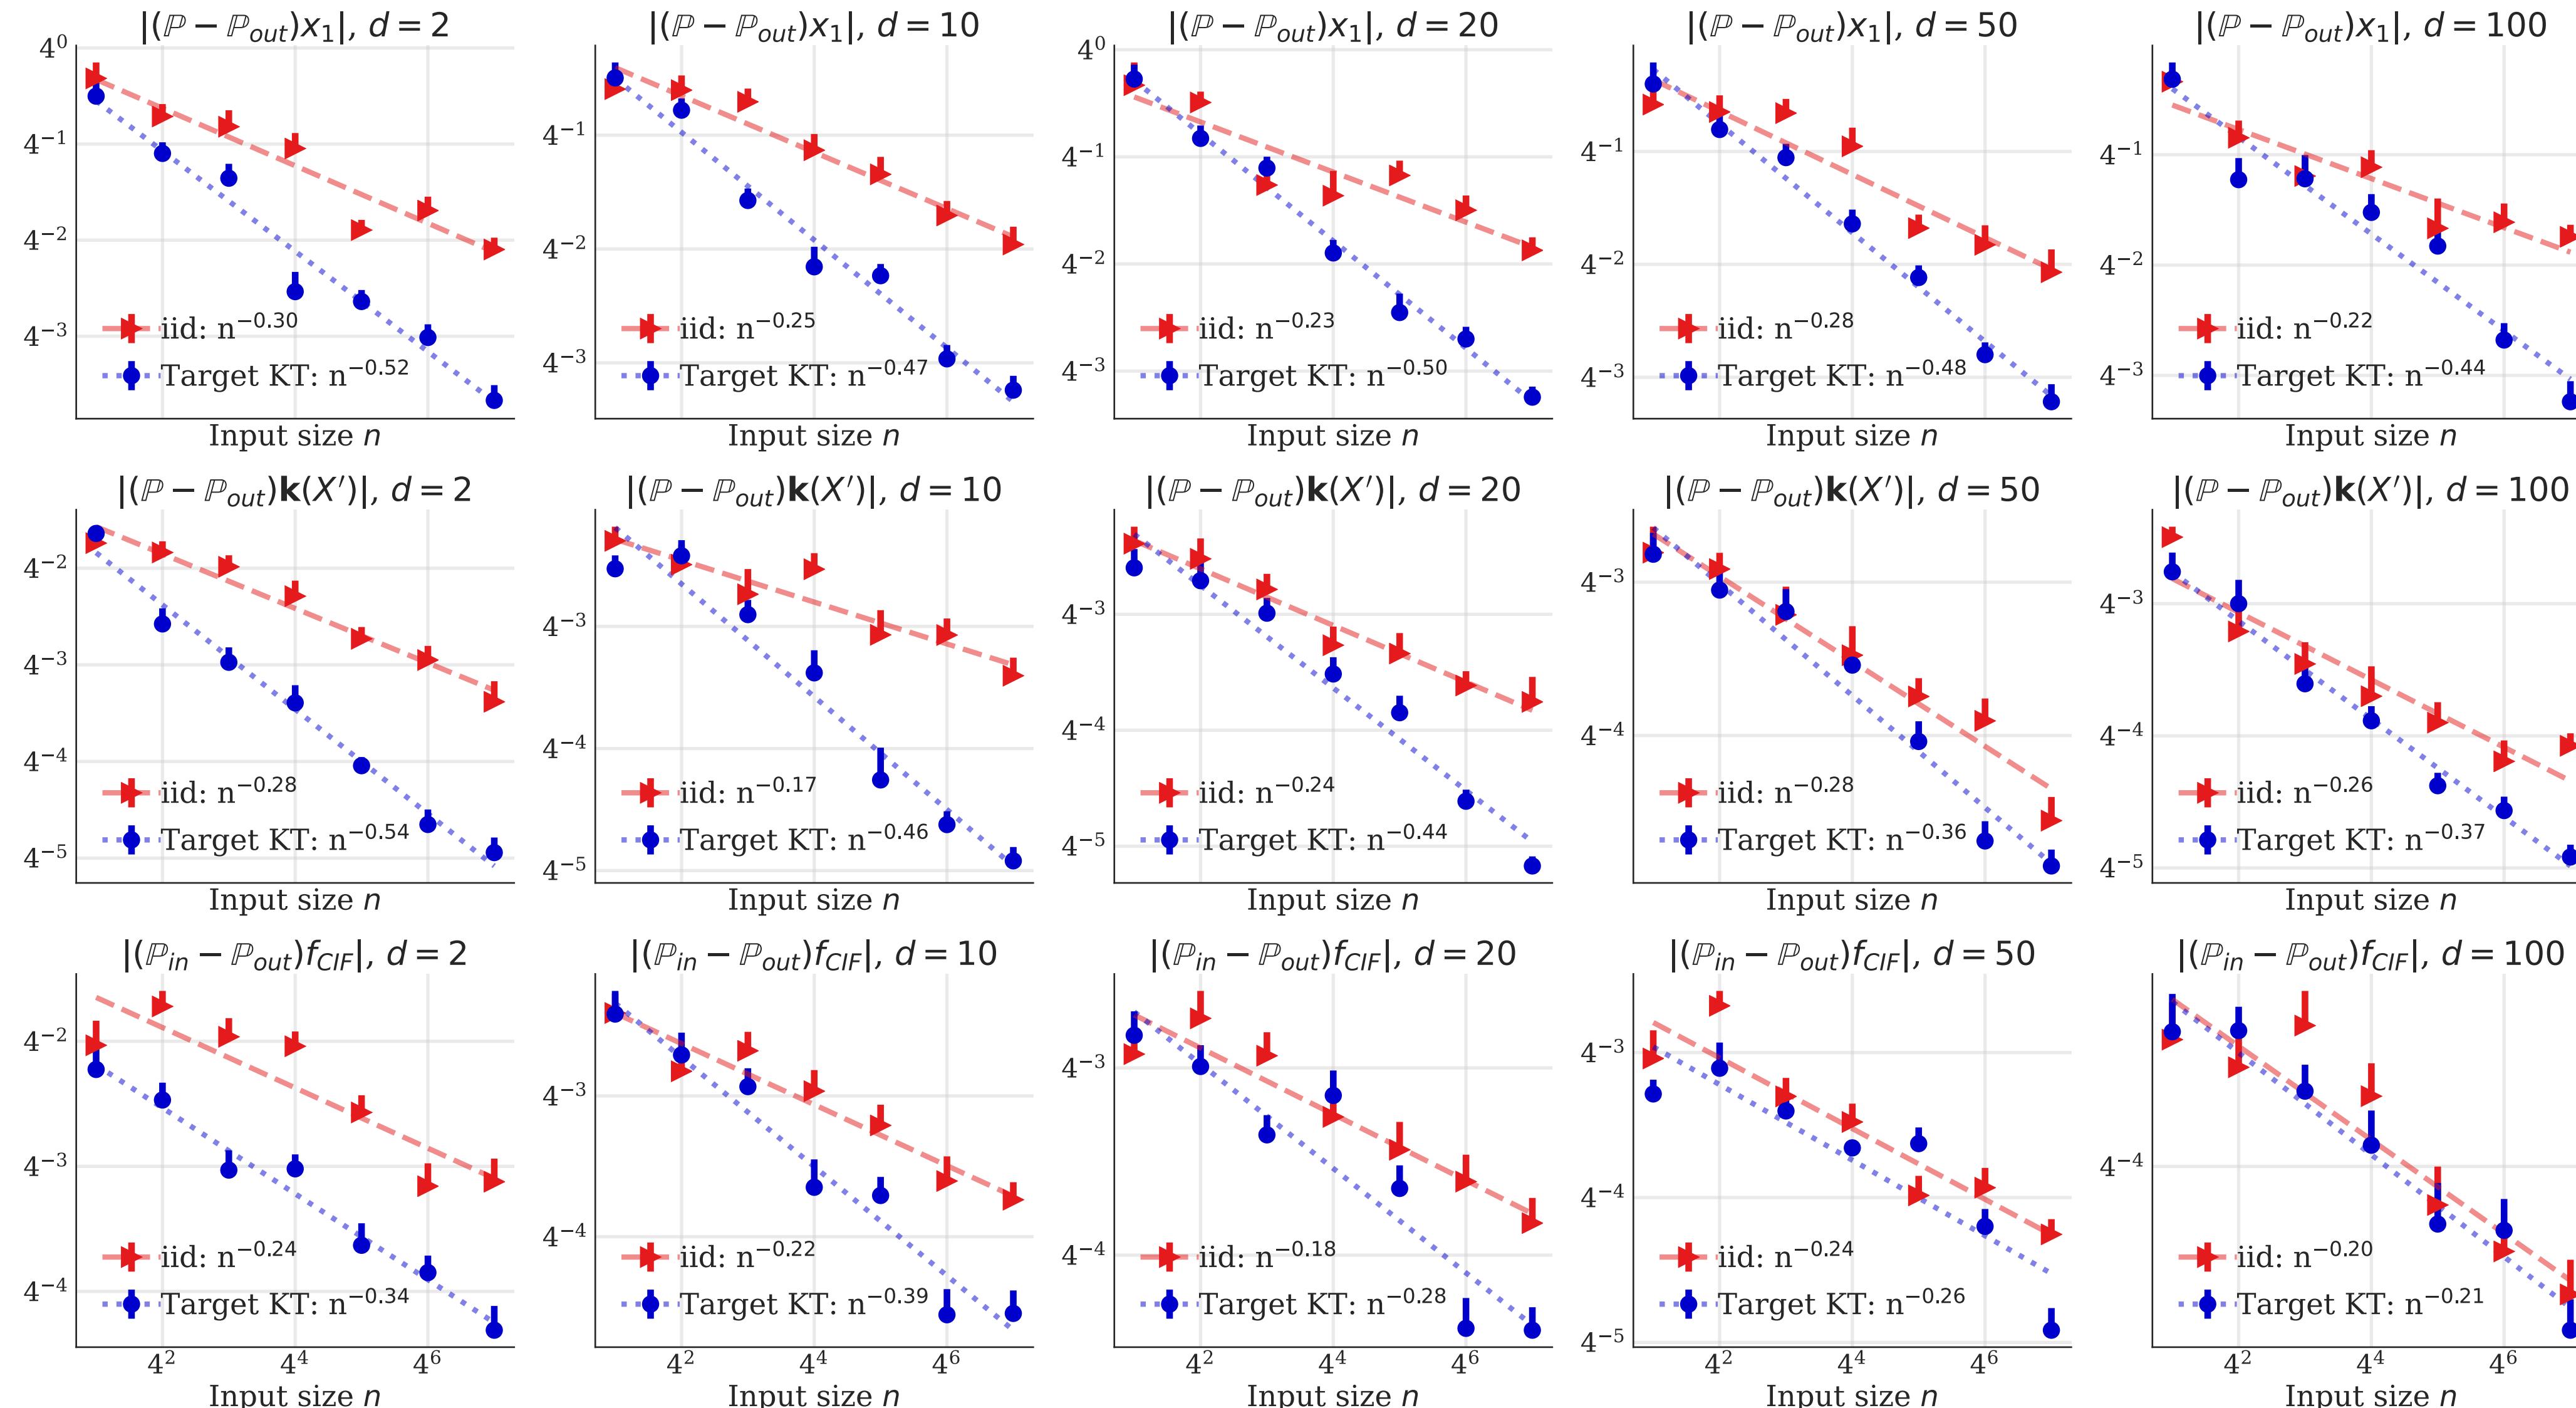
$$\widetilde{O}\left(\frac{n^{d/2m}}{\sqrt{n}}\right) \text{ for } m\text{-times differentiable kernels}$$

- We state explicit constants with dependence on kernel hyper-parameters in the paper

# **Single function experiments**

# Better error for functions inside and outside of RKHS

(Gaussian  $\mathbf{k}$  with  $\sigma^2 = 2d$  and standard Gaussian  $\mathbb{P}^\star$ )



# Contrast with related work

# Related work: $\sqrt{n}$ points with $n^{-1/4}$ MMD

- **Known guarantees *no better than Monte Carlo rate*:**

Standard thinning iid points [Tolstikhin-Sriperumbudur-Muandet, 2017]

Standard thinning geometrically ergodic MCMC [Dwivedi-Mackey 2021]

Kernel herding for infinite-dimensional kernels [Chen-Welling-Smola 2010, Lacoste-Julien-Lindsten-Bach 2015]

Stein Points MCMC [Chen-Barp-Briol-Gorham-Girolami-Mackey-Oates, 2019]

Greedy sign selection [Karnin-Liberty 2019]

- ***Unknown guarantees*:**

Support points [Mak-Joseph 2018]

Supersampling from a reservoir [Paige-Sejdinovic-Wood, 2016]:

# Related work: Better than Monte Carlo MMD

- **Finite-dimensional linear kernels:** Discrepancy construction [Harvey and Samadi, 2014]
- **Uniform  $\mathbb{P}^*$  on  $[0,1]^d$  (*bounded support*):**  
Quasi Monte Carlo [Hickernell 1998, Novak-Wozniakowski 2010],  
Haar thinning [Dwivedi-Feldheim-Gurel-Gurevich-Ramdas 2019]
- **$\mathbb{P}^*$  with *bounded support with known*  $\mathbb{P}^*k$ :**  
Bayesian quadrature [O'Hagan 1991]  
Bayes' Sard cubature [Karvonen et al. 2018]  
Determinantal point processes [Belhadji et al. 2020]
- **$(k, \mathbb{P}^*)$  with *known/bounded eigenfunctions*:**  
Determinantal point process kernel quadrature [Belhadji et al. 2019]  
Black-box importance sampling [Liu et al. 2018]

# Kernel thinning advantages

1.  $\sqrt{n}$  points with  $\widetilde{O}(n^{-1/2})$ -MMD error (iid sampling gives  $\Omega(n^{-1/4})$  error)
2. Valid for **non-uniform** target distributions with **unbounded support**
3. Valid for **infinite-dimensional** smooth/decaying kernels
4. Valid for **generic input points** including iid/MCMC/quadrature/herding with mild conditions
5. Requires **only kernel evaluations** to implement
6. **Matches MMD lower bounds** up to log factors
7. **Matches  $L^\infty$ -error lower bounds** up to log factors